

Operating Systems

Scheduling

Lecture Overview

- Scheduling

Last Week

Virtualisation

- Each process thinks it has its own computer
- Direct execution is faster
- Limited execution at key points to ensure OS control
- Hardware provides a lot of OS functionality
 - User vs Kernel
 - Clock
 - Register Saving

Scheduling

Who, when?

Scheduler Summary

- First Come, First Serve (FCFS)
- Shortest Job First
- Shortest Time to Completion First (STCF)
- Round Robin (RR)
- Multilevel Feedback Queue (MLFQ)

Scheduling vs Dispatching


Dispatcher

- Determines how to switch between processes
- User mode/kernel mode
- Saving PCBs...
- Mechanical (**How**)

Scheduler

- Determines which process to dispatch
- Conceptual (**When**)

Scheduling and States




Definitions

- **Workload:** The ‘**jobs**’ that need to be done
- **Job:** View as current CPU burst of a process
 - Process sometimes alternates between CPU & I/O
 - Process moves between ready and blocked queues
- **Scheduler:** Logic for the job choice
- **Metric:** Something to determine **good** from **bad**


Good Scheduling vs Bad Scheduling

Activity:


- Imagine you are running a hospital and you have a *operating room* (for surgery).
- You have a list of patients who need various operations...
- How do we determine the best way to use the *operating room*?
- What things do we want to **maximise** or **minimise**?




Some Answers




Some Answers: Turnaround Time




Some Answers: Response Time



Some Answers




Some Answers: Wait Time




Some Answers: Throughput




Some Answers: Resource Utilisation



Some Answers: Overhead (Switches)




Some Answers: Fairness




Scheduling Algorithms

Start with the basics...

GANTT Charts



Tables




Job	Arrival Time	Burst Time
Job	Arrival Time	Burst Time
A	0	8
B	5	7
C	10	6

Scheduling Assumptions

1. Jobs run for a fixed time
2. All jobs arrive simultaneously
3. All jobs are CPU-only
4. Run-time of each job is known


First Come, First Serve (FCFS)




First Come, First Serve (FCFS)

Metrics

- Turnaround Time $\text{completion_time} - \text{arrival_time}$
- Response Time




Job	Arrival Time	Burst Time
A	0	5
B	0	5
C	0	5




First Come, First Serve (FCFS)

Metrics

- Turnaround Time $\text{completion_time} - \text{arrival_time}$
- Response Time




Job	Arrival Time	Burst Time	Turnaround Time
A	0	5	5
B	0	5	10
C	0	5	15




First Come, First Serve (FCFS)

Metrics

- Turnaround Time $\text{completion_time} - \text{arrival_time}$
- Response Time



Job	Arrival Time	Burst Time	Turnaround Time
A	0	5	5
B	0	5	10
C	0	5	15




$$\frac{5 + 10 + 15}{3} = 10$$

Average


First Come, First Serve (FCFS)

Metrics

- Turnaround Time $\text{completion_time} - \text{arrival_time}$
- Response Time



Job	Arrival Time	Burst Time	Turnaround Time
A	0	20	20
B	0	4	24
C	0	4	28




$$\frac{20 + 24 + 28}{3} = 24$$

Average


First Come, First Serve (FCFS)

Metrics

- Turnaround Time $\text{completion_time} - \text{arrival_time}$
- Response Time



Job	Arrival Time	Burst Time	Turnaround Time
A	0	4	4
B	0	4	8
C	0	20	28



$$\frac{4 + 8 + 28}{3} = 16.6$$

Average

First Come, First Serve (FCFS)

Advantages

- Very simple to implement
- Has a ‘veneer’ of fairness


Problems

- Dependent on job order
- Big tasks can ‘starve’ small tasks



The Convoy Effect


The Convoy Effect




Shortest Job First (SJF)

Metrics

- Turnaround Time $\text{completion_time} - \text{arrival_time}$
- Response Time



Job	Arrival Time	Burst Time	Turnaround Time
A	0	20	20
B	0	4	24
C	0	4	28




$$\frac{4 + 8 + 28}{3} = 16.6$$

Average


Shortest Job First (SJF)

Metrics

- Turnaround Time $\text{completion_time} - \text{arrival_time}$
- Response Time



Job	Arrival Time	Burst Time	Turnaround Time
A	0	20	20
B	2	4	22
C	2	4	26




$$\frac{20 + 22 + 26}{3} = 22.6$$

Average

Scheduling Assumptions

1. ~~Jobs run for a fixed time~~
2. ~~All jobs arrive simultaneously~~
3. All jobs are CPU-only
4. Run-time of each job is known

Pre-emptive Scheduling




Task B

Task C



Non-Pre-emptive




Pre-emptive


Shortest Time to Completion First (STCF)

Metrics

- Turnaround Time $\text{completion_time} - \text{arrival_time}$
- **Response Time** $\text{start_time} - \text{arrival_time}$



Job	Arrival Time	Burst Time	Turnaround Time	Response Time
A	0	10	10	0
B	2	10	18	8
C	4	10	26	16




$$\frac{0 + 8 + 16}{3} = 8$$

Average


Round Robin (RR)

Metrics

- Turnaround Time $\text{completion_time} - \text{arrival_time}$
- Response Time $\text{start_time} - \text{arrival_time}$




Job	Arrival Time	Burst Time	Turnaround Time	Response Time
A	0	10	26	0
B	2	10	26	0
C	4	10	26	0



$$\frac{0 + 0 + 0}{3} = 0$$


Average

Round Robin (RR)




Metrics

- Turnaround Time
- Response Time



Turnaround Time = completion_time – arrival_time
Response Time = start_time – arrival_time

Job	Arrival Time	Burst Time	Turnaround Time	Response Time
A	0	10	26	0
B	2	10	26	0
C	4	10	26	0



$$\frac{0 + 0 + 0}{3} = 0$$

Average


Input/Output Blocking

What are you waiting for?


Scheduling Assumptions

1. ~~Jobs run for a fixed time~~
2. ~~All jobs arrive simultaneously~~
3. ~~All jobs are CPU only~~
4. Run-time of each job is known


I/O Blocking



I/O Blocking



I/O Blocking



Don't waste CPU

Scheduling Assumptions


1. ~~Jobs run for a fixed time~~
2. ~~All jobs arrive simultaneously~~
3. ~~All jobs are CPU only~~
4. ~~Run time of each job is known~~

Multi-Level Feedback Queue (MLFQ)

Situation

- Jobs can be of varying length
- Jobs can arrive at any time
- We don't know when they will finish
- Different jobs have different aims
 - Interactive Jobs: **Response Time**
 - Batch Jobs: **Turnaround Time**

Priority (Fixed)



Priority (Fixed)



Round Robin

Priority (Fixed)



Big Job

Lower Priority


It will never run...

Priority (Dynamic)



Runs too long?

Drop priority



Tiny I/O Jobs



Logic:


If a job takes a lot of CPU (i.e., everything you give it), reduce its priority...

If a job uses only a little CPU, keep its priority

Priority (Dynamic)

**If a job can only ‘lose priority’,
won’t it eventually get ignored?**


Multi-level Feedback QUEUE




Priority (Evil)

Under this system, how can we be evil?


Priority (Dynamic)



Queue 2



Queue 1



Queue 0




MLFQ

Rules:

- Run high priority jobs first (round robin for ties)
- You start high priority
- If you run for too long, your priority drops
- Every so often, put jobs back into the top priority

Priority (Evil)



Allotment

Leave a tiny gap 😊

If we 90% of our allotment, our priority will never fall...

MU HA HA HA HA HA HA HA

MLFQ

Rules:

- Run high priority jobs first (round robin for ties)
- You start high priority
- If you run for **use a full allotment**, your priority drops
- Every so often, put jobs back into the top priority

MLFQ

Design Decisions:

- How many queues?
- How long is an allotment?
 - Are allotments the same for each queue level?
- How often to refresh priorities?


Lottery Scheduling

Choose processes ‘at random’

- Processes get tickets, and you choose a random ticket
- Priorities => More Tickets

Some weird thoughts

- Ticket trading?
- Ticket inflation?
- Ticket currencies?



Completely Fair Scheduler (CFS)

Premise:

- All tasks have their **vruntime** measured
 - We can mess with this by changing the weighting of the increments of runtime
- Smallest **vruntime** tasks are run with priority


- Each task will run for a bit over a **target latency** period
- Each task gets a share depending on its weight (so no starvation)
- Works using a Red/Black tree

Run in Linux since 2007

Completely Fair Scheduler (CFS)

Why?

- Priority is important to know
- Finding the highest priority needs to be low complexity ($O(\log(n))$)



Completely Fair Scheduler (CFS)

Example

- Task 1: Video Rendering
 - Long run-time, CPU intensive
 - Non-interactive
- Task 2: Word Processor
 - Intermittent, CPU bursts
 - Interactive

Result

Task 2 accumulates little **vruntime**. When it wakes (user input) it gets run immediately.

Completely Fair Scheduler (CFS)

nice(1)

The ‘**nice**’ can be thought of as a ‘soft’ priority.

- Changing a program’s **nice** value changes the weighting for CFS.





The ‘**nice**’ values correspond to the weightings so you get a proportion of CPU time
(lower nice = more CPU)

Multi-core CPUs

Linux

- One red-black tree per core
But...

How to fairly distribute tasks between cores?



Summary

Scheduling:

- First Come, First Serve FCFS (also FIFO)
 - Round-Robin RR
 - Shortest Job First SJF
 - Multi-level Feedback Queues MLFQ

Questions?

