

# Operating Systems

Concurrency

# Lecture Overview

- Locks
- Lock Implementation
  - Data Structures
- Condition Variables

# Last Week


## **Memory Virtualisation**

- Page Swapping


# Atomicity and Concurrency

Things happening at the same time or not at the same time

# Multiple CPUs



# Multiple CPUs



One Process, Two Threads


# The problem of Atomicity

Consider:

```
incrementBalance ()  
{  
    balance = balance + 1;  
}
```

What this translates to:

```
mov 0x123, %eax  
add %0x1, %eax  
mov %eax, 0x123
```




# The problem of Atomicity

Consider:

```
incrementBalance ()  
{  
    balance = balance + 1;  
}
```

What this translates to:

```
mov 0x123, %eax  
add %0x1, %eax  
mov %eax, 0x123
```



# Activity

You have two threads.

x, i, o all start = 0

What will be printed out?

$x = x + 1$

$i = i + 1$

Print(x, i)


$x = x + 1$

$o = o + 1$

Print(x, o)


**Red, Orange, Yellow, Green, Blue, Purple = GOOD**

**Red, Green, Anything = BAD**



Red, Orange, Yellow, Green, Blue, Purple = GOOD

**Red, Green, Anything = BAD**



# Race Conditions...

This is a single 'logical' step

```
mov 0x123, %eax  
add %0x1, %eax  
mov %eax, 0x123
```

Mess with 0x123

The atomic steps are no longer atomic

# The Conch Shell

## **Lord of the Flies**

- Only the person with the conch shell can speak



# Kicking the Can Down the Road

## What if?

```
incrementBalance()  
{  
    while (editing)  
    {  
        // Do nothing  
    }  
    editing = true;  
    balance = balance + 1;  
    editing = false;  
}
```

## Explanation:

- Two processes call increment balance...
- One sets editing to **true**
- The other waits until editing becomes **false**

# Kicking the Can Down the Road

## What if?

```
incrementBalance ()  
{  
    while (someone_has_conch)  
    {  
        // Do nothing  
    }  
    someone_has_conch = true;  
    balance = balance + 1;  
    someone_has_conch = false;  
}
```

## Explanation:

- Two processes call increment balance...
- One sets someone\_has\_conch to **true**
- The other waits until someone\_has\_conch becomes **false**

What is the problem here?

# Kicking the Can Down the Road

## What if?

```
incrementBalance ()  
{  
    while (someone_has_conch)  
    {  
        // Do nothing  
    }  
    someone_has_conch = true;  
    balance = balance + 1;  
    someone_has_conch = false;  
}
```

Process 1

Check  
'conch'

Has  
'conch'

No  
'conch'

Process 2

Check  
'conch'

Has  
'conch'

Has  
'conch'

# Kicking the Can Down the Road

## What if?

```
incrementBalance ()  
{  
    while (someone_has_conch)  
    {  
        // Do nothing  
    }  
    someone_has_conch = true;  
    balance = balance + 1;  
    someone_has_conch = false;  
}
```



Process 1

Check  
'conch'

Has  
'conch'

No  
'conch'

Process 2

Check  
'conch'

Has  
'conch'


Has  
'conch'

# Conch is not enough...

## Requirement:

- Check availability
- Take

} Must be atomic



# Mutexes in C

How to write the code...


# Mutex

## Mutual Exclusion in C

```
// Setup  
pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;  
...  
// Use  
pthread_mutex_lock(&lock);  
//Atomic code  
balance = balance + 1;  
pthread_mutex_unlock(&lock);
```

We can have many 😊

# Example: Linked List



# Example: Linked List

```
typedef struct __node_t
{
    int key;
    struct __node_t *next;
} node_t;
```

```
typedef struct __list_t
{
    node_t * head;
} list_t;
```


```
void ListInit(list_t *L)
{
    L->head = NULL;
}
```

# Example: Linked List

```
void ListInsert(list_t *L, int key)
{
    node_t *new_node = malloc(sizeof(node_t));
    assert(new_node);
    new_node->key = key;
    new_node->next = L->head;
    L->head = new_node;
}

int ListLookup(list_t *L, int key)
{
    node_t *temp = L->head;
    while (temp)
    {
        if (temp->key == key)
            return 1;
        temp = temp->next;
    }
    return 0;
}
```

How to break this?




# Example: Linked List

```
void ListInsert(list_t *L, int key)
{
    node_t *new_node = malloc(sizeof(node_t));
    assert(new_node);
    new_node->key = key;
    new_node->next = L->head;
    L->head = new_node;
}
```

Thread A

```
new_node->key = key;
new_node->next = L->head;
```

```
L->head = new_node;
```



Thread B

```
new_node->key = key;
new_node->next = L->head;
L->head = new_node;
```

# Example: Linked List

```
typedef struct __node_t
{
    int key;
    struct __node_t *next;
} node_t;
```

```
typedef struct __list_t
{
    node_t * head;
    pthread_mutex_t lock;
} list_t;
```

```
void ListInit(list_t *L)
{
    L->head = NULL;
    pthread_mutex_init(&L->lock, NULL);
}
```

## Initialisation

# Example: Linked List: Approach #1

```
void ListInsert(list_t *L, int key)                                int ListLookup(list_t *L, int key)
{
    node_t *new_node = malloc(sizeof(node_t));
    assert(new_node);
    new_node->key = key;
    new_node->next = L->head;
    L->head = new_node;
}

atomic action {
```

```
    node_t *temp = L->head;
    while (temp)
    {
        if (temp->key == key)
            return 1;
        temp = temp->next;
    }
    return 0;
}
```

Where are the critical sections?

# Example: Linked List: Approach #1


```
void ListInsert(list_t *L, int key)
{
    pthread_mutex_lock(&L->lock);

    node_t *new_node = malloc(sizeof(node_t));
    assert(new_node);

    new_node->key = key;

    new_node->next = L->head;
    L->head = new_node;

    pthread_mutex_unlock(&L->lock);
}
```




```
int ListLookup(list_t *L, int key)
{
    pthread_mutex_lock(&L->lock);

    node_t *temp = L->head;
    while (temp)
    {
        if (temp->key == key)
            return 1;

        temp = temp->next;
    }

    pthread_mutex_unlock(&L->lock);
    return 0;
}
```




How to break this?

# Example: Linked List: Approach #2

```
void ListInsert(list_t *L, int key)
{
    node_t *new_node = malloc(sizeof(node_t));
    assert(new_node);
    new_node->key = key;
    new_node->next = L->head;
    L->head = new_node;
}
```

atomic  
action



```
int ListLookup(list_t *L, int key)
{
    node_t *temp = L->head;
    while (temp)
    {
        if (temp->key == key)
            return 1;
        temp = temp->next;
    }
    return 0;
}
```

Where are the critical sections?

# Example: Linked List: Approach #2

```
void ListInsert(list_t *L, int key)
{
    node_t *new_node = malloc(sizeof(node_t));
    assert(new_node);
    new_node->key = key;
    pthread_mutex_lock(&L->lock);
    new_node->next = L->head;
    L->head = new_node;
    pthread_mutex_unlock(&L->lock);
}

int ListLookup(list_t *L, int key)
{
    pthread_mutex_lock(&L->lock);
    node_t *temp = L->head;
    while (temp)
    {
        if (temp->key == key)
        {
            pthread_mutex_unlock(&L->lock);
            return 1;
        }
        temp = temp->next;
    }
    pthread_mutex_unlock(&L->lock);
    return 0;
}
```

What is the smallest atomic action?

# Lock Implementation

How to make it actually work?

# Implementation

## Design Goals

- Correct
  - Mutual Exclusion
  - No Deadlock
  - Bounded (starvation free)
- Fast & Cheap
- Fair

## Notes

- Software solution appears impossible...

# Bad Idea #1: Disable interrupts

## Rationale

- No interrupts prevents the dispatcher from running another thread.

## Problems

- Lock the CPU (evil process)
- Only works in a uni-processor environment

```
void lock(locT *l)
{
    disableInterrupts();
}
```

```
void lock(locT *l)
{
    enableInterrupts();
}
```

This is sometimes a viable tactic in **kernel** programming (embedded systems)


# Bad Idea #2: Load & Store

```
typedef struct __lock_t {  
    int flag;  
} lock_t;
```

```
void init(lock_t * mutex)  
{  
    mutex->flag = 0;  
}
```

“Spin locks” are  
**really** inefficient

```
void lock(lock_t * mutex)  
{  
    while (mutex->flag == 1)  
    {  
        mutex->flag = 1;  
    }  
}  
  
void unlock(lock_t * mutex)  
{  
    mutex->flag = 0;  
}
```



# Bad Idea #2: Load & Store

```
void lock(lock_t * mutex)
{
    while (mutex->flag == 1)

    mutex->flag = 1;
}
```

```
while (mutex->flag == 1)
{ }
```

```
mutex->flag = 1;
```

```
void lock(lock_t * mutex)
{
    while (mutex->flag == 1)

    mutex->flag = 1;
}
```

```
while (mutex->flag == 1)
{ }
mutex->flag = 1;
```

# Idea #3: Peterson's Algorithm

```
int turn = 0;
boolean flag[2];

void init()
{
    flag[0] = flag[1] = false;
    turn = 0;
}

void lock()
{
    flag[self] = true;
    turn = 1 - self;
    while(flag[1-self] && (turn == 1-self))
    {
    }

    void unlock()
    {
        flag[self] = false;
    }
}
```

There are two flags? One per thread

# Idea #3: Peterson's Algorithm

## Intuition

- Do stuff if:
  - Other thread isn't interested
  - Other thread gave you permission
- Progress
  - No infinite waiting (one thread or other must pass)
- Starvation
  - Each thread only waits once

## But...

- Cache?
- The caches across CPUs may not be consistent
- Instructions may not actually be run in order (modern processors)
- Write to memory in execution order (cache activity, optimisations)


# Memory Barrier Instruction

**Idea:**

- Memory accesses may not cross the barrier

**C:**


```
_sync_synchronize();
```



# XCHG: Atomic Exchange “Test and Set”

```
// xchg(int *addr, int newval)
// return what was pointed to by addr
// at the same time, store newval inaddr

int xchg(int * addr, int newval)
{
    int old = *addr;
    *addr = newval;
    return old;
}
```



## Rationale?

- You can set something and ensure that the ‘thing’ changed? Or didn’t change?

# XCHG: Atomic Exchange “Test and Set”

```
void lock(lock_t * lock)
{
    while (xchg(&lock->flag, 1) == 1)
    {
    }
}
```

```
void unlock(lock_t * lock)
{
    lock->flag = 0;
}
```

## Logic:

- Set lock to one...
- But if it is already one...
- Do nothing and try again

# Compare and Swap

```
int CAS(int * addr, int expected, int new)
{
    int actual = *addr;
    if (actual == expected)
        *addr = new;
    return actual;
}

void lock(lock_t * lock)
{
    while (CAS(&lock->flag, 0, 1) == 1)
    {
    }
}
```

## Subtle Difference:

- What is it? What should it be?
- If it is what it should be update it...
- Either way return what it is?

# Implementation

## Design Goals

- Correct
  - Mutual Exclusion
  - No Deadlock
  - Bounded (starvation free)
- Fast & Cheap
- Fair  ?

## Notes:

- Hardware support


# Ticket Locks

## Issue:

- Spin-locks are arbitrary?

Another  
thread gets  
in...


Mutex  
becomes  
free



# Ticket Locks

## Issue:

- Spin-locks are arbitrary?



# Ticket Locks


## Issue:

- Spin-locks are arbitrary?

## Idea:

- Have an incremented value (atomically) while returning the old value
- Spin while ‘not your turn’
- Increment turn when done...

```
int FetchAndAdd(int *ptr)
{
    int old = *ptr;
    *ptr = old + 1;
    return old;
}
```



# Ticket Locks

A lock()

B lock()

C lock()

0
1
2
3
4
5
6
7

# Ticket Locks

A lock() → Ticket 0

B lock() → Ticket 1

C lock() → Ticket 2



# Ticket Locks

A lock() → Ticket 0

B lock() → Ticket 1

C lock() → Ticket 2

A runs




# Ticket Locks

A lock() → Ticket 0

B lock() → Ticket 1

C lock() → Ticket 2

A runs → turn = 1



# Ticket Locks


A lock() → Ticket 0

B lock() → Ticket 1

C lock() → Ticket 2

A runs → turn = 1

B runs



# Ticket Locks

A lock() → Ticket 0

B lock() → Ticket 1

C lock() → Ticket 2

A runs → turn = 1

B runs

A lock() → Ticket 3



# Ticket Locks

A lock() → Ticket 0


B lock() → Ticket 1

C lock() → Ticket 2

A runs → turn = 1

B runs → turn = 2

A lock() → Ticket 3



# Ticket Locks

A lock() → Ticket 0

B lock() → Ticket 1


C lock() → Ticket 2

A runs → turn = 1

B runs → turn = 2

A lock() → Ticket 3

C runs



# Ticket Locks

A lock() → Ticket 0

B lock() → Ticket 1

C lock() → Ticket 2


A runs → turn = 1

B runs → turn = 2

A lock() → Ticket 3

C runs

B lock() → Ticket 4




# Spinlocks

## Fast If:

- Many CPUs
- Locks held for short time
- (no context switch required)

## Slow If:



- One CPU
- Locks held for long time
- (spinning is just bad)



Why run B, C, D if they are  
waiting for A?

What does it achieve?

# Yielding



```
typedef struct __lock_t
{
    int ticket;
    int turn;
} lock_t;

int FetchAndAdd(int *ptr)
{
    int old = *ptr;
    *ptr = old + 1;
    return old;
}
```

```
void acquire(lock_t *lock)
{
    int myturn = FetchAndAdd(&lock->ticket);
    while (lock->turn != myturn)
        yield();
}

void release(lock_t *lock)
{
    FetchAndAdd(&lock->turn);
}
```


# Evaluating Lock Designs

## Fairness:

- Is request order the same as receipt order?

## Performance:

- Low contention (lock mostly there)
- High contention (lock mostly in use)



This is effectively  
scheduling

# Spinlock Performance Comparison

## CPU Cost:

- Without yield:  $O(\text{threads} * \text{time\_slice})$
- With yield:  $O(\text{threads} * \text{context\_switch})$
- Yielding is still expensive

# Implementation: Block when Waiting


## Concept:

- If a thread is ‘waiting’, put it on the blocked queue
- Scheduler only runs ‘ready’ processes

Here the OS scheduler doesn’t need to be involved (i.e., separation of concerns)


# Example

Running: A  
Ready: B, C, D  
Blocked:




# Example

Running:            B  
Ready:             C, D, A  
Blocked:




# Example

Running:	C
Ready:	D, A
Blocked:	B




# Example

Running:	D
Ready:	A, C
Blocked:	B




# Example

Running:	A
Ready:	C
Blocked:	B, D




# Example

Running:	C
Ready:	A
Blocked:	B, D




# Example

Running: A  
Ready: C, B  
Blocked: D



# Example



B is probably feeling a bit hard done by...

# Code: Block When Waiting

```
typedef struct {
    bool flag;
    bool guard;
    queue_t *q;
} lock_t;

void acquire(lock_t * l) {
    while (TestAndSet(&l->guard, 1) == 1);
    if (l->flag) {
        queue_add(l->q, gettid());
        l->guard = 0;
        park(); // Blocked
    } else {
        l->flag = 1; // Lock acquired
        l->guard = 0;
    }
}

void lock_init(lock_t *m)
{
    m->flag = 0;
    m->guard = 0;
    queue_init(m->q);
}

void release(lock_t * l) {
    while (TestAndSet(&l->guard, 1) == 1);
    if (queue_empty(l->q)) {
        l->flag = 0;
    } else {
        unpark(queue_remove(l->q));
    }
    l->guard = 0;
}
```

# Code: Block When Waiting

```
typedef struct {
    bool flag;
    bool guard;
    queue_t *q;
} lock_t;

void acquire(lock_t * l) {
    while (TestAndSet(&l->guard, 1) == 1);
    if (l->flag) {
        queue_add(l->q, gettid());
        l->guard = 0;
        park(); // Blocked
    } else {
        l->flag = 1; // Lock acquired
        l->guard = 0;
    }
}

void lock_init(lock_t *m)
{
    m->flag = 0;
    m->guard = 0;
    queue_init(m->q);
}

void release(lock_t * l) {
    while (TestAndSet(&l->guard, 1) == 1);
    if (queue_empty(l->q)) {
        l->flag = 0;
    } else {
        unpark(queue_remove(l->q));
    }
    l->guard = 0;
}
```

We get stuck by the guard...

# Code: Block When Waiting

```
typedef struct {
    bool flag;
    bool guard;
    queue_t *q;
} lock_t;

void acquire(lock_t * l) {
    while (TestAndSet(&l->guard, 1) == 1);
    if (l->flag) {
        queue_add(l->q, gettid());
        l->guard = 0;
        park(); // Blocked
    } else {
        l->flag = 1; // Lock acquired
        l->guard = 0;
    }
}

void lock_init(lock_t *m)
{
    m->flag = 0;
    m->guard = 0;
    queue_init(m->q);
}

void release(lock_t * l) {
    while (TestAndSet(&l->guard, 1) == 1);
    if (queue_empty(l->q)) {
        l->flag = 0;
    } else {
        unpark(queue_remove(l->q));
    }
    l->guard = 0;
}
```

Flag should normally be zero to start, so we set the lock and reset the guard

# Code: Block When Waiting

```
typedef struct {
    bool flag;
    bool guard;
    queue_t *q;
} lock_t;

void acquire(lock_t * l) {
    while (TestAndSet(&l->guard, 1) == 1);
    if (l->flag) {
        queue_add(l->q, gettid());
        l->guard = 0;
        park(); // Blocked
    } else {
        l->flag = 1; // Lock acquired
        l->guard = 0;
    }
}

void lock_init(lock_t *m)
{
    m->flag = 0;
    m->guard = 0;
    queue_init(m->q);
}

void release(lock_t * l) {
    while (TestAndSet(&l->guard, 1) == 1);
    if (queue_empty(l->q)) {
        l->flag = 0;
    } else {
        unpark(queue_remove(l->q));
    }
    l->guard = 0;
}
```

The second process will get added to a queue, reset the guard and then get parked()

# Code: Block When Waiting

```
typedef struct {
    bool flag;
    bool guard;
    queue_t *q;
} lock_t;

void acquire(lock_t * l) {
    while (TestAndSet(&l->guard, 1) == 1);
    if (l->flag) {
        queue_add(l->q, gettid());
        l->guard = 0;
        park(); // Blocked
    } else {
        l->flag = 1; // Lock acquired
        l->guard = 0;
    }
}

void lock_init(lock_t *m)
{
    m->flag = 0;
    m->guard = 0;
    queue_init(m->q);
}

void release(lock_t * l) {
    while (TestAndSet(&l->guard, 1) == 1);
    if (queue_empty(l->q)) {
        l->flag = 0;
    } else {
        unpark(queue_remove(l->q));
    }
    l->guard = 0;
}
```

Again, we're going to worry about the guard

# Code: Block When Waiting

```
typedef struct {
    bool flag;
    bool guard;
    queue_t *q;
} lock_t;

void acquire(lock_t * l) {
    while (TestAndSet(&l->guard, 1) == 1);
    if (l->flag) {
        queue_add(l->q, gettid());
        l->guard = 0;
        park(); // Blocked
    } else {
        l->flag = 1; // Lock acquired
        l->guard = 0;
    }
}

void lock_init(lock_t *m)
{
    m->flag = 0;
    m->guard = 0;
    queue_init(m->q);
}

void release(lock_t * l) {
    while (TestAndSet(&l->guard, 1) == 1);
    if (queue_empty(l->q)) {
        l->flag = 0;
    } else {
        unpark(queue_remove(l->q));
    }
    l->guard = 0;
}
```

Here we unpark the 2<sup>nd</sup> process  
(the head of the queue)

# Code: Block When Waiting

```
void acquire(lock_t * l) {
    while (TestAndSet(&l->guard, 1) == 1);
    if (l->flag) {
        queue_add(l->q, gettid());
        l->guard = 0;
        park(); // Blocked
    } else {
        l->flag = 1; // Lock acquired
        l->guard = 0;
    }
}

void release(lock_t * l) {
    while (TestAndSet(&l->guard, 1) == 1);
    if (queue_empty(l->q)) {
        l->flag = 0;
    } else {
        unpark(queue_remove(l->q));
    }
    l->guard = 0;
}
```

Some Questions:

**The Guard?**

**Spinning Guard?**

**Unpark doesn't reset Flag?**

**Race condition?**

# Code: Block When Waiting

```
void acquire(lock_t * l) {  
    while (TestAndSet(&l->guard, 1) == 1);  
    if (l->flag) {  
        queue_add(l->q, gettid());  
        l->guard = 0;  
        park(); // Blocked  
    } else {  
        l->flag = 1; // Lock acquired  
        l->guard = 0;  
    }  
}  
  
void release(lock_t * l) {  
    while (TestAndSet(&l->guard, 1) == 1);  
    if (queue_empty(l->q)) {  
        l->flag = 0;  
    } else {  
        unpark(queue_remove(l->q));  
    }  
    l->guard = 0;  
}
```

Some Questions:

**The Guard?**

**Spinning Guard?**

**Unpark doesn't reset Flag?**

**Race condition?**

The guard acts to separate flag setting and queuing behaviours?

# Code: Block When Waiting

```
void acquire(lock_t * l) {  
    while (TestAndSet(&l->guard, 1) == 1);  
    if (l->flag) {  
        queue_add(l->q, gettid());  
        l->guard = 0;  
        park(); // Blocked  
    } else {  
        l->flag = 1; // Lock acquired  
        l->guard = 0;  
    }  
}  
  
void release(lock_t * l) {  
    while (TestAndSet(&l->guard, 1) == 1);  
    if (queue_empty(l->q)) {  
        l->flag = 0;  
    } else {  
        unpark(queue_remove(l->q));  
    }  
    l->guard = 0;  
}
```

Some Questions:

**The Guard?**

**Spinning Guard?**

**Unpark doesn't reset Flag?**

**Race condition?**

The Guard spins minimally...

# Code: Block When Waiting

```
void acquire(lock_t * l) {  
    while (TestAndSet(&l->guard, 1) == 1);  
    if (l->flag) {  
        queue_add(l->q, gettid());  
        l->guard = 0;  
        park(); // Blocked  
    } else {  
        l->flag = 1; // Lock acquired  
        l->guard = 0;  
    }  
}  
  
void release(lock_t * l) {  
    while (TestAndSet(&l->guard, 1) == 1);  
    if (queue_empty(l->q)) {  
        l->flag = 0;  
    } else {  
        unpark(queue_remove(l->q));  
    }  
    l->guard = 0;  
}
```

Some Questions:

**The Guard?**

**Spinning Guard?**

**Unpark doesn't reset Flag?**

**Race condition?**

Allows new process ‘handover’

# Race Condition

```
if (l->flag) {  
    queue_add(l->q, gettid());  
    l->guard = 0;
```

```
park();
```

I was told to wakeup just before  
I went to sleep (instead of after)  
Now I must sleep FOREVER

```
while (TestAndSet(&l->guard, 1) == 1);  
if (queue_empty(l->q)) {  
    l->flag = 0;  
} else {  
    unpark(queue_remove(l->q));  
}  
l->guard = 0;
```

# Code: Block When Waiting: Fix

```
typedef struct {
    bool flag;
    bool guard;
    queue_t *q;
} lock_t;

void acquire(lock_t * l) {
    while (TestAndSet(&l->guard, 1) == 1);
    if (l->flag) {
        queue_add(l->q, gettid());
        setpark(); // Notify of plan
        l->guard = 0;
        park(); // Blocked
    } else {
        l->flag = 1; // Lock acquired
        l->guard = 0;
    }
}

void lock_init(lock_t *m)
{
    m->flag = 0;
    m->guard = 0;
    queue_init(m->q);
}

void release(lock_t * l) {
    while (TestAndSet(&l->guard, 1) == 1);
    if (queue_empty(l->q)) {
        l->flag = 0;
    } else {
        unpark(queue_remove(l->q));
    }
    l->guard = 0;
}
```


If I am about to sleep and someone  
'wakes' me, I won't go to sleep anymore.

# Spin-Locking vs Blocking

How much does a context switch cost (C)?


## Uniprocessor:

- Waiting process is scheduled, process holding lock isn't
- Waiting process should always relinquish processor
- Associate queue of waiters with each lock



## Multiprocessor:

- Waiting process is scheduled, process holding lock could be?
- Spin or block depending on time before lock release
- Lock released **quickly**: spin
- Lock released **slowly**: block





# Spin-Locking vs Blocking

How much does a context switch cost (C)?

## Question:

- We will wait  $t$  for the lock to be released
- If  $C > t$ :
  - Spin lock
- If  $C < t$ :
  - Block



# Futex

## Overview:

- futex is a linux-based implementation of mutex-style behaviour
  - Queue: Kernel
  - Atomic Integer: User space
- Main takeaway:
  - If the threads don't contend, the kernel system calls are never invoked ☺

`futex_wait(address, expected)`

- Puts the calling thread to sleep
- If the value at address is not expected, return call immediately



`futex_wake(address)`

- Wake one thread waiting in the queue

# Condition Variables

Threads aren't independent


# Relay Race



# Logic

## The Handover:

- **Wait** until they are done
- **Signal** when you are done



# Wait!

## Function Signature:

```
pthread_cond_wait(&cond, &lock)
```

**cond:** pthread\_cond\_t

**lock:** pthread\_mutex\_t

**storage\_t:** mailbox

# Signal!

## Function Signature:

```
pthread_cond_signal(&cond)
```


**cond:** pthread\_cond\_t

**lock:** pthread\_mutex\_t

**storage\_t:** mailbox

Does nothing if no one  
is actually waiting

# In Pictures



# In Code

```
mutex_lock (&lock)
// do stuff
done = true
cond_signal(&cond)
mutex_unlock(&lock)
```

```
mutex_lock (&lock)
if (!done)
    cond_wait(&cond, &lock)
//do stuff
mutex_unlock(&lock)
```

# In Code

```
mutex_lock(&lock)  
// do stuff  
done = true  
cond_signal(&cond)  
mutex_unlock(&lock)
```

Sneak in here

```
mutex_lock(&lock)  
if (!done)  
    cond_wait(&cond, &lock)  
//do stuff  
mutex_unlock(&lock)
```

# In Code

```
mutex_lock (&lock)  
// do stuff  
done = true  
cond_signal (&cond)  
mutex_unlock (&lock)
```

```
mutex_lock (&lock)  
if (!done)  
cond_wait (&cond, &lock)  
//do stuff  
mutex_unlock (&lock)
```

What will we wait for?

# Pipes

Ceci n'est pas une pipe


# Example

```
ls | wc
```


# Pipes






# Pipes



# Pipes



# Pipes: Buffer



read





end



NOW WE WAIT...

# Pipes: Buffer



# Pipes: Producers and Consumers

**Producer:** Generate Data (pipe writer)


**Consumer:** Accept data and process it (pipe reader)

## Examples:

- Server, Client (webservers)
- Pipes


# Example: Code

```
void * producer(void * arg) {  
    for (int i=0; i<loops; i++) {  
        mutex_lock(&m);  
        while (numfull == max)  
            cond_wait(&cond, &m);  
        do_fill(i);  
        cond_signal(&cond);  
        mutex_unlock(&m);  
    }  
}  
  
void * consumer(void * arg) {  
    while (1) {  
        mutex_lock(&m);  
        while (numfull == 0)  
            cond_wait(&cond, &m);  
        int tmp = do_get();  
        cond_signal(&cond);  
        mutex_unlock(&m);  
        printf("%d\n", tmp);  
    }  
}
```




# Example: Code

```
void * producer(void * arg) {  
    for (int i=0; i<loops; i++) {  
        mutex_lock(&m);  
        while (numfull == max)  
            cond_wait(&cond, &m);  
        do_fill(i);  
        cond_signal(&cond);  
        mutex_unlock(&m);  
    }  
}  
  
void * consumer(void * arg) {  
    while (1) {  
        mutex_lock(&m);  
        while (numfull == 0)  
            cond_wait(&cond, &m);  
        int tmp = do_get();  
        cond_signal(&cond);  
        mutex_unlock(&m);  
        printf("%d\n", tmp);  
    }  
}
```




# Example: Code

```
void * producer(void * arg) {  
    for (int i=0; i<loops; i++) {  
        mutex_lock(&m);  
        while (numfull == max)  
            cond_wait(&cond, &m);  
        do_fill(i);  
        cond_signal(&cond);  
        mutex_unlock(&m);  
    }  
}  
  
void * consumer(void * arg) {  
    while (1) {  
        mutex_lock(&m);  
        while (numfull == 0)  
            cond_wait(&cond, &m);  
        int tmp = do_get();  
        cond_signal(&cond);  
        mutex_unlock(&m);  
        printf("%d\n", tmp);  
    }  
}
```




# Example: Code


```
void * producer(void * arg) {  
    for (int i=0; i<loops; i++) {  
        mutex_lock(&m);  
        while (numfull == max)  
            cond_wait(&cond, &m);  
        do_fill(i);  
        cond_signal(&cond);  
        mutex_unlock(&m);  
    }  
}  
  
void * consumer(void * arg) {  
    while (1) {  
        mutex_lock(&m);  
        while (numfull == 0)  
            cond_wait(&cond, &m);  
        int tmp = do_get();  
        cond_signal(&cond);  
        mutex_unlock(&m);  
        printf("%d\n", tmp);  
    }  
}
```



# Why is this useful?



# Why is this useful?



# Condition Variables: Fairness

Which thread to wake?

**signal** (cond\_t)

- Wake one thread
- Else do nothing

**wait** (cond\_t, mutex\_t)

- Assumes lock held prior to call
- Sleeps and releases locks
- Reacquires lock on waking

**broadcast** (cond\_t)

- Wake all threads
- Else do nothing

# Example: Code

```
void * producer(void * arg) {  
    for (int i=0; i<loops; i++) {  
        mutex_lock(&m);  
        while (numfull == max)  
            cond_wait(&full, &m);  
        do_fill(i);  
        cond_signal(&empty);  
        mutex_unlock(&m);  
    }  
}  
  
void * consumer(void * arg) {  
    while (1) {  
        mutex_lock(&m);  
        while (numfull == 0)  
            cond_wait(&empty, &m);  
        int tmp = do_get();  
        cond_signal(&full);  
        mutex_unlock(&m);  
        printf("%d\n", tmp);  
    }  
}
```

Maybe have two condition variables?

# Condition Variables Summary

## Notes:

- Always do wait/signal while holding a mutex
- Always have a separate concept of ‘state’
- Whenever a thread wakes, recheck

# Summary

## Concurrency

- Locks
- Lock Implementation
  - Data Structures
- Condition Variables

# Questions?

