

Operating Systems

I/O, Hard Drives

Lecture Overview

- Canonical Device
- Direct Memory Access
- Hard Drives
- Solid State Drives

Last Week

Concurrency

- Semaphores

I/O Devices

What to do?


Motivation

What purpose is a computer that doesn't interact with the outside world?

A bit of history

Original UI


- Physical switches Input
- Lights Output




A bit of history

Punch Cards


- Detect holes (physically)
- Print holes (physically)



Architecture



Architecture




Device Controllers

- Local Buffer Storage
- Set of Special Purpose Registers

Canonical Device

Common Registers

- Writing to a register may cause something to happen
- Reading from a register may cause something to happen
- Both can cause something to happen
- What you read isn't what you wrote



Example


```
while (STATUS == BUSY)
    ; // spin

Write data to DATA register

Write command to COMMAND register

while (STATUS == BUSY)
    ; // spin
```

This is a very
simplified example



Example

```
while (STATUS == BUSY)
```

```
    ; // spin
```

Write data to DATA register

Write command to COMMAND register

```
while (STATUS == BUSY)
```

```
    ; // spin
```

Burning CPU
(probably Kernel mode)

CPU

A

DEVICE


Example

```
while (STATUS == BUSY)  
    ; // spin
```

Write data to DATA register

Write command to COMMAND register

```
while (STATUS == BUSY)  
    ; // spin
```




Example

```
while (STATUS == BUSY)  
    ; // spin
```

Write data to DATA register

Write command to COMMAND register

```
while (STATUS == BUSY)  
    ; // spin
```



Example


```
while (STATUS == BUSY)  
    ; // spin
```

Write data to DATA register

Write command to COMMAND register

```
while (STATUS == BUSY)  
    ; // spin
```

Writing
COMMAND



Example

```
while (STATUS == BUSY)
```

```
    ; // spin
```

Write data to DATA register

Write command to COMMAND register

```
while (STATUS == BUSY)
```

```
    ; // spin
```

Do I/O Stuff




Example

```
while (STATUS == BUSY)  
    ; // spin
```

Write data to DATA register

Write command to COMMAND register

```
while (STATUS == BUSY)  
    ; // spin
```



Example

```
while (STATUS == BUSY)
```

wait for interrupt

Write data to DATA register


Write command to COMMAND register

```
while (STATUS == BUSY)
```

wait for interrupt

Feels pretty inefficient...

Couldn't we instead of **BUSY** waiting, perhaps we can perform an **interrupt**



Example

```
while (STATUS == BUSY)
```

wait for interrupt

Write data to DATA register



Write command to COMMAND register

```
while (STATUS == BUSY)
```

wait for interrupt

Feels pretty inefficient...

Couldn't we instead of **BUSY** waiting, perhaps we can perform an **interrupt**





Example

```
while (STATUS == BUSY)  
    wait for interrupt  
Write data to DATA register  
Write command to COMMAND register  
while (STATUS == BUSY)  
    wait for interrupt
```

Feels pretty inefficient...

Couldn't we instead of **BUSY** waiting, perhaps we can perform an **interrupt**



Interrupts vs Polling

Polling

- Wastes CPU time
- Can work on fast devices

Interrupts

- Interrupt overheads
- Can result in live-lock (interrupt hell)
 - Better to ignore interrupts while some make progress
- Interrupt coalescing

Hybrid

- Spin, then use interrupts

Connecting I/O Devices

How is it done?

Protocol Variants

Port-mapped I/O


- Have an address space reserved for devices
- Use specialised CPU instruction (in, out) to read/write to this space
- Separates devices from normal memory
- Often one-byte at a time
- Examples: x86

Memory-mapped I/O

- Device registers are simply mapped into RAM
- OS reads these locations like any other location
- Simplified implementation
- Should not be **cached**

Protocol Variants

Port-mapped I/O



Memory-mapped I/O



Direct Memory Access

What would you say, you do around here?

Problem

I/O Device

- I would like to copy some data into memory

CPU

- That is my job!

Problem


I/O Device

- I would like to copy some data into memory

CPU

- That is my job!

Copy this
byte



Problem


I/O Device

- I would like to copy some data into memory

CPU

- That is my job!

Copy this
byte




Extremely
inefficient
CPU usage


In Pictures: Programmed I/O

Process

- CPU copies a byte
- CPU copies a byte
- ...



In Pictures: Direct Memory Access



Updating our model

PIO

```
while (STATUS == BUSY)
```

wait for interrupt

Write data to DATA register

Write command to COMMAND register

```
while (STATUS == BUSY)
```

wait for interrupt

DMA

```
while (STATUS == BUSY)
```

wait for interrupt

~~*Write data to DATA register*~~

Write command to COMMAND register

```
while (STATUS == BUSY)
```

wait for interrupt

Device Drivers

Writing lots of code


Device Driver

os

- Kernel stuff

Device

- Works a particular way



Device Driver

Device Drivers


- 70% of linux code is **drivers**

Mac


- Integrated HW (i.e., no issues)
- 3rd party driver instructions

Windows

- Historically separate driver CD's



Kernel Device Structure



Device Drivers

What is it?

- Device-specific code in the kernel that interacts directly with the device hardware
 - (Translation of the generic interface)
- Allows a generic interface

Device Drivers

Top Half


- kernel's interface to the **device driver** (standardised interface)
 - Linux "everything is a file"
- read() / write()
- open() / close()
- ioctl() system call (Linux)
 - Special device-specific configuration supported by the

Bottom Half

- Interrupt service routines
- DMA operations


Lifecycle of an I/O request

We request some I/O
use in our program




Lifecycle of an I/O request

We transfer the data,
and return completion
(or **error**)




Lifecycle of an I/O request

We're done, either we have received **input** or pushed **output**




Lifecycle of an I/O request

Now we need to talk to
the device driver
Block the process




Lifecycle of an I/O request




Lifecycle of an I/O request


time to send the commands to the physical hardware



Lifecycle of an I/O request




Lifecycle of an I/O request




Lifecycle of an I/O request

register that the interrupt arrived queue? unblock device




Lifecycle of an I/O request

do work (copying etc)
inform kernel that the actual
work has been done




Lifecycle of an I/O request




back to transfer data!

Lifecycle of an I/O request



Everything is a file

```
FILE fd = fopen("/dev/something", "rw");  
for (int i = 0; i < 10; i++) {  
    fprintf(fd, "Count %d\n", i);  
}  
close(fd);
```



Everything is a file...

Block Devices: Drives

- Access blocks of data
- Commands:
 - open/read/write/**seek**
- Raw I/O or file-system access
- Memory mapped file access possible



Everything is a file...

```
emper@Dooly:/dev$ ls -l
total 0
crw-r--r-- 1 root root      10, 235 Sep  1 16:55 autofs
drwxr-xr-x 2 root root      580 Sep  1 16:55 block
drwxr-xr-x 2 root root      100 Sep  1 16:55 bsg
crw-rw---- 1 root disk     10, 234 Sep  1 16:55 btrfs-control
drwxr-xr-x 3 root root      60 Sep  1 16:55 bus
drwxr-xr-x 2 root root     2780 Sep  1 16:55 char
crw----- 1 root root      5,  1 Sep  1 16:55 console
lrwxrwxrwx 1 root root      11 Sep  1 16:55 core -> /proc/kcore
crw----- 1 root root     10, 125 Sep  1 16:55 cpu_dma_latency
crw----- 1 root root     10, 203 Sep  1 16:55 cuse
drwxr-xr-x 6 root root      120 Sep  1 16:55 disk
drwxr-xr-x 3 root root      100 Sep  1 16:55 dri
crw-rw-rw- 1 root root     10, 127 Sep  1 16:55 dxg
lrwxrwxrwx 1 root root      13 Sep  1 16:55 fd -> /proc/self/fd
crw-rw-rw- 1 root root      1,  7 Sep  1 16:55 full
crw-rw-rw- 1 root root     10, 229 Sep  1 16:55 fuse
drwxr-xr-x 2 root root      0 Sep  1 16:55 hugepages
crw----- 1 root root    229,  0 Sep  1 16:55 hvc0
crw--w---- 1 root tty     229,  1 Sep  1 16:55 hvc1
crw----- 1 root root    229,  2 Sep  1 16:55 hvc2
crw----- 1 root root    229,  3 Sep  1 16:55 hvc3
crw----- 1 root root    229,  4 Sep  1 16:55 hvc4
crw----- 1 root root    229,  5 Sep  1 16:55 hvc5
crw----- 1 root root    229,  6 Sep  1 16:55 hvc6
crw----- 1 root root    229,  7 Sep  1 16:55 hvc7
```

d:	directory
l:	link
c:	character device file
b:	block device file
p:	named pipe
s:	socket
tty:	terminal

Everything is a file

Character Devices: keyboards, mice, serial ports

- Single characters at a time
- Commands:
 - `get()`, `put()`
- Libraries layered on top allow line editing



Everything is a file

Network Devices: Ethernet, Wireless, Bluetooth

- Different enough to have their own interface
- Use sockets() (see CAN)
- Functions
 - select()
- Usage: pipes, FIFOs, streams, queues, mailboxes



Timing Paradigms

Blocking Interface: “Wait”

- When request data, ‘sleep’ process until data ready
- When write data, ‘sleep’ until device is ready for data

Non-blocking Interface: “Don’t wait”

- Returns quickly from read or write request with count of bytes transferred
- Read may return nothing, write may return nothing

Asynchronous Interface: “Tell me later”

- When request data, take pointer to user’s buffer, return immediately; later kernel fills buffer and notifies user
- When send data, take pointer to user’s buffer, return immediately; later kernel takes data and notifies user

I/O Device Types

Summary

- Different I/O
 - Speeds (0.1 bytes/s => Gbytes/s)
 - Access patterns?
 - Access timing? (blocking, non-blocking, asynchronous)
 - Notification mechanisms (interrupts vs polling)
 - Types (block, character, network)

Hard Disks

Huzzahs

Tape Drives

Used for storage


- Need to wind tape to find correct section of ‘information’ to load into memory.
- Incredibly slow!
- Incredibly dense!



Hard Disk Drive

Used for storage


- Wheel allows quick access to different locations on the disk
- Pretty quick?
- Pretty dense?



Hard Disk Drive


Contain:

- Tracks
- Sectors (512B or 4096B)
- Cylinders Many heads...
- Platters
- Spindle




Reading/Writing

Spindle... spins



Reading/Writing



Some questions:

- How does the disk head know where it is?
 - Platters misaligned
 - Tracks not perfectly concentric
- High precision manufacturing
 - Some contain helium??

Reading/Write

Magnetic disk

- Pass a field over a region to change its state (lasts a long time).


Drive Servo Systems

- Need to work out where you are (lets you choose which sector to read/write to
 - Notches (old)
 - Timing information within a disk

Seek, Rotate, Transfer

Seek Cost:

- Depends on the cylinder distance
 - Not purely linear
 - Must accelerate/coast/slow/settle
- Entire seek often takes several milliseconds (4-10ms)
- Average seek distance $\sim 1/3$ of max seek distance



Seek, Rotate, Transfer

Rotate:

- Depends on the rotations per minute (RPM)
 - 3600
 - 5400
 - 7200
 - 15000


7200 RPM

⇒ 1 minute / 7200

⇒ 8.3 ms / rotation

Average rotation

$$= 8.3/2 = 4.15\text{ms}$$




Seek, Rotate, Transfer

Transfer

- Pretty fast
 - RPM (fast = more)
 - Data density (more = more)
 - Request size (max = more)
- Typically, GB/s

$$4096 \text{ B} * 1\text{GB/s} = 4 \text{ ms transfer}$$



Disk Comparison

	Cheetah	Barracuda
Capacity	300 GB	1 TB
RPM	15000	7200
Av Seek (ms)	4	9
Max Transfer (MB/s)	125	105
Platters	4	4
Cache MB	16	32

What is the throughput?

Cheetah: 125 MB/s
Barracuda: 105 MB/s

Disk Comparison: Calculations

	Cheetah	Barracuda
RPM	15000	7200
Av Seek (ms)	4	9
Max Transfer (MB/s)	125	105

How long would an average **random** 16 KB read take with the Cheetah?

Seek: 4ms

Disk Comparison: Calculations

	Cheetah	Barracuda
RPM	15000	7200
Av Seek (ms)	4	9
Max Transfer (MB/s)	125	105

How long would an average **random** 16 KB read take with the Cheetah?

$$\frac{1}{2} \times \frac{1}{15000} \times \frac{60}{1} \times \frac{1000}{1}$$

Unit Conversions

Seek: 4ms

2ms

Disk Comparison: Calculations

	Cheetah	Barracuda
RPM	15000	7200
Av Seek (ms)	4	9
Max Transfer (MB/s)	125	105

How long would an average **random** 16 KB read take with the Cheetah?

$$\frac{16}{125} \times \frac{1}{1024} \times \frac{1000}{1}$$

0.125ms

Seek: 4ms
Rotation: 2ms

Unit Conversions

Disk Comparison: Calculations

	Cheetah	Barracuda
RPM	15000	7200
Av Seek (ms)	4	9
Max Transfer (MB/s)	125	105

How long would an average **random** 16 KB read take with the Cheetah?

Seek	Rotate	Transfer	Seek:	4ms
4	2	0.125	Rotation:	2ms
			Transfer:	0.125ms

$$= 6.125$$

Buffering

Internal Memory

- 2-16 MB cache
- Read contents of entire track into memory during rotational delay
 - Cache leverages spatial locality
 - “Read ahead”

Write caching with volatile memory

- Immediate reporting
- Problem: You claim to have written when you haven’t (power failure = bad)

Tagged Command Queuing

- Have multiple outstanding requests (SATA... 16?)
- Disk can re-order requests

More Disk ‘Smarts’

Depends upon what the external abstraction is:

- Disks may re-order operations
- Makes reasoning about their performance very difficult
- Operating systems may not have any useful optimisations available
 - SCSI devices do not expose internals – just a linear space of blocks

Solid State Disks

Flash!

Solid State Disks (SSD)

No moving parts

- Flash memory
- Written in pages (not bytes)
- Not byte-level addressable
 - Read whole page
 - Modify
 - Write whole page
 - Can only write to erased blocks



SSD – Brief history

Battery-backed DRAM (1995)

- Battery needed to avoid data loss

NAND Multi-level Cell (2 or 3 bit cell) **flash memory** (2009)



1991

The logic: NAND Flash Memory


Memory consists of cells

- Originally 1-bit but can now be multi-bit

Memory is stored in **pages**

Pages are grouped into **blocks**

A page




A block

The logic: NAND Flash Memory

Pages are initially (empty – this usually actually means all 1's)

When you need a page you find one and write new memory into it.

A page




A block

The logic: NAND Flash Memory

If all the **pages** are full... you
can't overwrite them

Many pages




A block

The logic: NAND Flash Memory

Instead, you would need to
‘empty’ an entire **block**, then
write the new **page**.

New page




A block

The logic: NAND Flash Memory

You can't grab bytes from within a page.

You need to take the whole page.




A block

How can this possibly be good?

Where are the problems?


What can we do to fix these?

Copy and Paste



Old page is read, copied to a new page with modifications... then marked **invalid**


Garbage Collection



Garbage Collection

Once a block is mostly invalid, it can be scrubbed.

Deletion




TRIM Command

- Delete just means... mark a page as **invalid**


Flash Memory (1)

Current implies 1




Flash Memory (0)

No Current implies 0




Flash Memory (Writing zero)

No Current implies 0



Flash Memory (1)

Current implies 1



Flash Memory: Problems


Erasing does damage to the cell itself.

- ⇒ Each cell has a limited lifecycle (~10K)
- ⇒ minimise its use

SSD Architecture

Write 4KB Page ~200 μ s – 1.7 ms

- Only empty pages
- Erase ~ 1.5ms



New Bottlenecks

Hard Drives: Slow

⇒ Use SATA bus (600 MB/s)

SSD Drives: Fast

⇒ Use PCIe bus (7-14 GB/s)


M.2 Slots added

Software Changes needed too:

- AHCI (Advance Host Controller Interface)
 - Single command queue, 32 entries
- NVMe
 - Reduced latency
 - Massive parallelism

Wear Leveling


If memory blocks can wear out:



This is bad

Wear Leveling

If memory blocks can wear out:




This is better

Wear Leveling

If memory blocks can wear out:

But...

- 10K writes
- Static data is also bad... (can get stuck)
- Some data shuffling occurs



This is better

Wear Leveling

Layer indirection

- Maintain a Flash Translation Layer (FTL in SSD)
- Map virtual block numbers (SO) to physical page numbers (flash memory controller uses these)

The memory can move without the OS knowing

Copy on Write

- Do not overwrite a page when OS updates its data
- Instead write a new version to a free page.
- Update FTL mapping

SSD Summary

Pros (vs Hard Drives)


- Faster (lower latency, not seeking/rotating)
- No moving parts (less fragile)

Cons

- Assymetric performance (read/write/erase)
- Complicated
 - Wear Leveling
 - Garbage collections
- Lifetime (pretty minor)

Schedulers

Where to put it?



First Come, First Serve (FCFC)

You work this out...

Shortest Positioning Time First (SPTF)

Greedy Algorithm

- Choose the shortest one first

Implementation

OS


- Shortest Seek Time First
 - Starves distant requests

SCAN

Elevator Algorithm


- Go up and down sweeping and servicing requests as you pass them
 - Sorts by cylinder (ignores rotation delays)

Circular Scan



LOOK

Like SCAN, but reverse when you reach the last request.



I/O Device Summary


Summary

- Overlap I/O with CPU where possible
 - Interrupts and DMA are your friend
- Storage devices pretend to be ‘one big block’ (even when they are not)
- Never do random access I/O unless you really need to (i.e., linked lists are bad)
- Scheduling can be effective for otherwise slow devices

Filesystems

How?

From Storage to File Systems



An interesting issue

Windows:

C:\\

Linux

\

Summary

- Canonical Device
- Direct Memory Access
- Hard Drives
- Solid State Drives

Questions?

