

Operating Systems

Log-structured File System

Overview

- LFS: Log-Structured File System

Some Observations

- Memory gets bigger => Cache gets bigger => Read less frequent
- Hard Drives (density) gets better
- Hard Drive motors (seek, rotate) are mostly stagnant
- FFS: inode, data, bitmap, bitmap
- RAID is not part of the picture

Summary

Write >>> Read

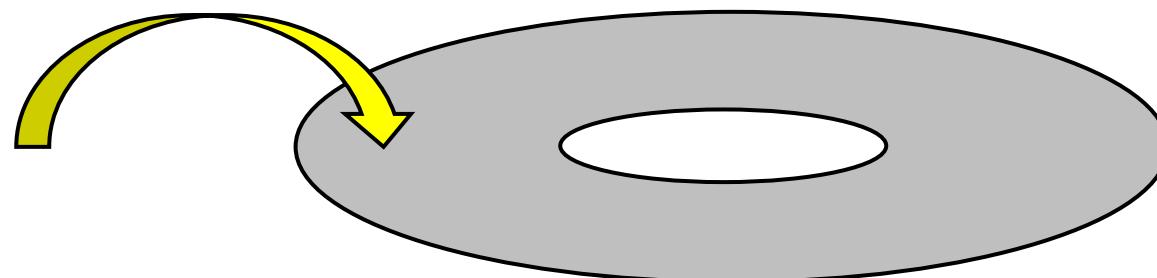
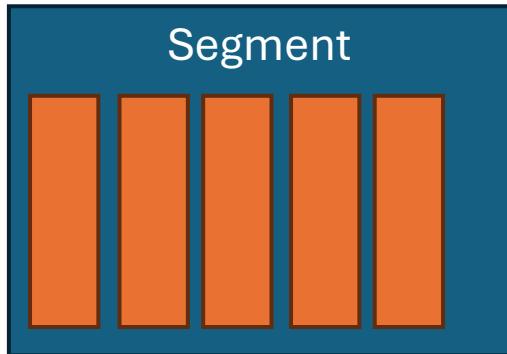
Ideas

Principles

- We want to minimize ‘writes’
- We want to write ‘slabs’ of data (not random access)

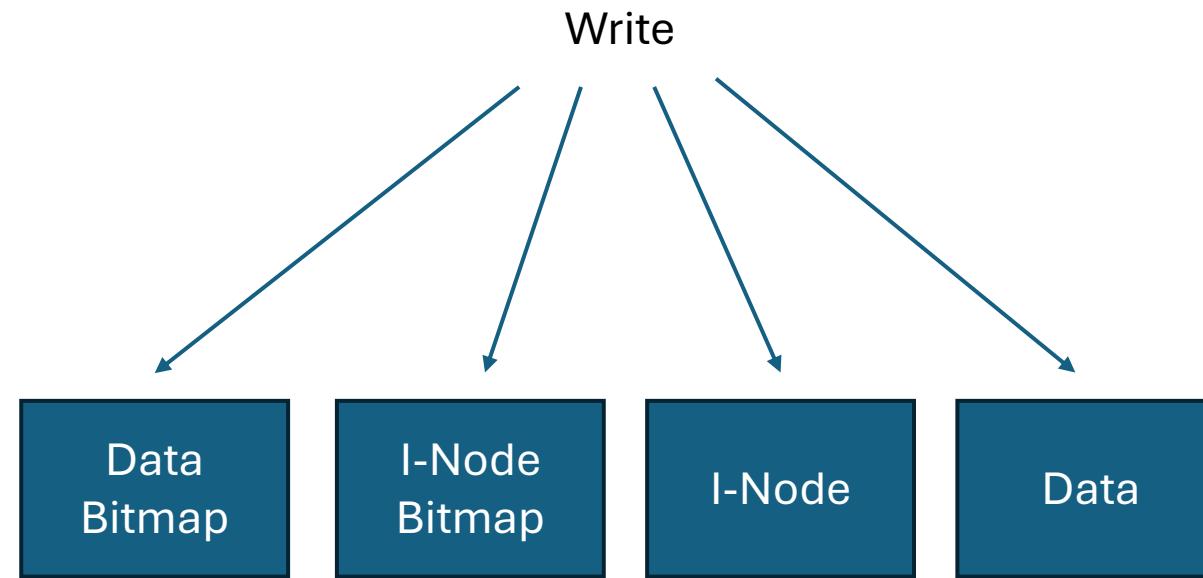
Theory

- Buffer all the reads into a segment
- Write the segment
- Write to a new location



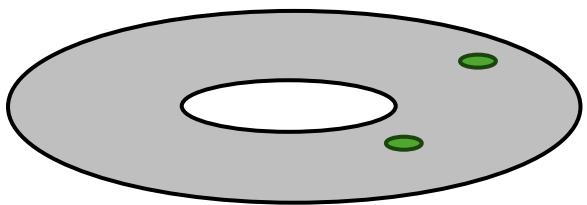
Problems

#1: Writing everything together isn't 'free'. We need to think about this.

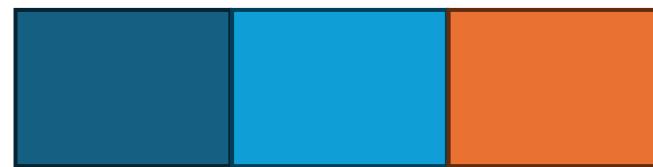


Problems

#1: Writing everything together isn't 'free'. We need to think about this.



Sequential

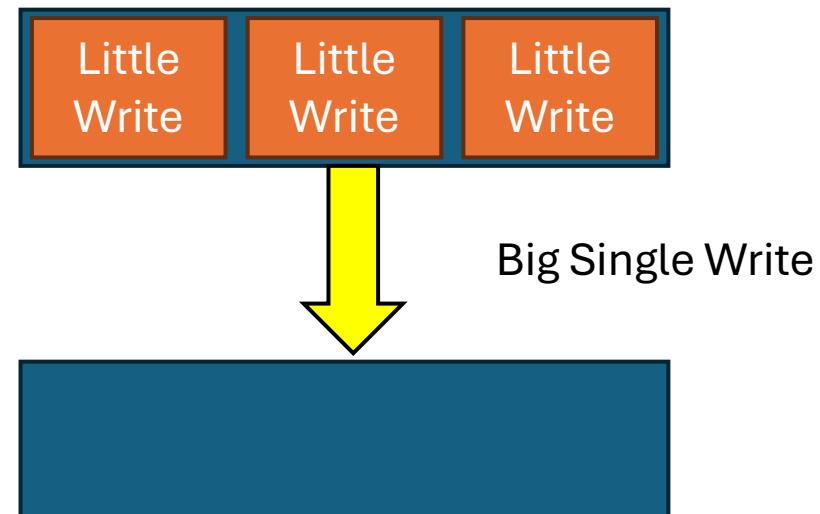
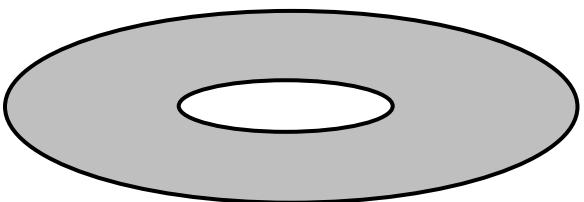


Contiguous



Problems

#1: Writing everything together isn't 'free'. We need to think about this.



Problems

#2: How much to buffer?

Activity

	Big Buffer	Small Buffer
Good		
Bad		

Thoughts

	Big Buffer	Small Buffer
Good	<ul style="list-style-type: none">• Less write operations• Better if seek/rotation time is high	<ul style="list-style-type: none">• Less ‘waiting’ for memory writes
Bad	<ul style="list-style-type: none">• Could lose a lot of data if it isn’t written	<ul style="list-style-type: none">• More write operations• Worse if seek/rotation time is high

Maths

$$Rate_{effective} = \frac{Data}{T_{write}} = \frac{Data}{T_{position} + \frac{Data}{Rate_{peak}}}$$

Seek/Rotation Time

Maths

We want:

$$\text{Rate}_{\text{effective}} = 0.9 \times \text{Rate}_{\text{peak}}$$

$$\text{Rate}_{\text{effective}} = \frac{\text{Data}}{T_{\text{write}}} = \frac{\text{Data}}{T_{\text{position}} + \frac{\text{Data}}{\text{Rate}_{\text{peak}}}}$$

Maths

We want:

$$\text{Rate}_{\text{effective}} = 0.9 \times \text{Rate}_{\text{peak}}$$

$$\text{Rate}_{\text{effective}} = \frac{\text{Data}}{\text{T}_{\text{position}} + \frac{\text{Data}}{\text{Rate}_{\text{peak}}}}$$

Maths

We want:

$$\text{Rate}_{\text{effective}} = 0.9 \times \text{Rate}_{\text{peak}}$$

$$\text{Rate}_{\text{effective}} = \frac{\text{Data}}{T_{\text{position}} + \frac{\text{Data}}{\text{Rate}_{\text{peak}}}}$$

$$0.9 * \text{Rate}_{\text{peak}} = \frac{\text{Data}}{T_{\text{position}} + \frac{\text{Data}}{\text{Rate}_{\text{peak}}}}$$

Maths

We want:

$$\text{Rate}_{\text{effective}} = 0.9 \times \text{Rate}_{\text{peak}}$$

$$0.9 * \text{Rate}_{\text{peak}} = \frac{\text{Data}}{T_{\text{position}} + \frac{\text{Data}}{\text{Rate}_{\text{peak}}}}$$

$$\text{Data} = 0.9 * \text{Rate}_{\text{peak}} * \left(T_{\text{position}} + \frac{\text{Data}}{\text{Rate}_{\text{peak}}} \right)$$

Maths

We want:

$$\text{Rate}_{\text{effective}} = 0.9 \times \text{Rate}_{\text{peak}}$$

$$\text{Data} = 0.9 * \text{Rate}_{\text{peak}} * \left(T_{\text{position}} + \frac{\text{Data}}{\text{Rate}_{\text{peak}}} \right)$$

$$\text{Data} = 0.9 * \text{Rate}_{\text{peak}} * T_{\text{position}} + 0.9 * \text{Rate}_{\text{peak}} * \frac{\text{Data}}{\text{Rate}_{\text{peak}}}$$

Maths

We want:

$$\text{Rate}_{\text{effective}} = 0.9 \times \text{Rate}_{\text{peak}}$$

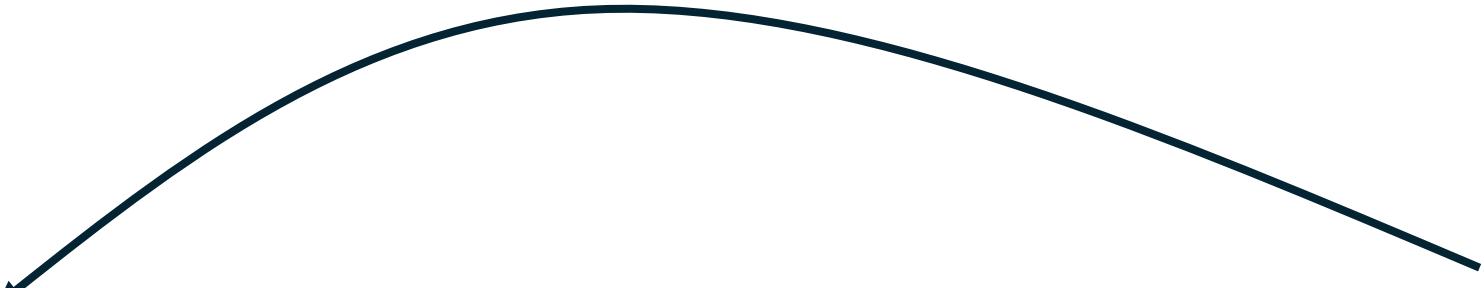
$$\text{Data} = 0.9 * \text{Rate}_{\text{peak}} * T_{\text{position}} + 0.9 * \cancel{\text{Rate}_{\text{peak}}} * \frac{\text{Data}}{\cancel{\text{Rate}_{\text{peak}}}}$$

$$\text{Data} = 0.9 * \text{Rate}_{\text{peak}} * T_{\text{position}} + 0.9 * \text{Data}$$

Maths

We want:

$$\text{Rate}_{\text{effective}} = 0.9 \times \text{Rate}_{\text{peak}}$$


$$Data = 0.9 * Rate_{peak} * T_{position} + 0.9 * Data$$

$$Data - 0.9 * Data = 0.9 * Rate_{peak} * T_{position}$$

Maths

We want:

$$\text{Rate}_{\text{effective}} = 0.9 \times \text{Rate}_{\text{peak}}$$

$$Data - 0.9 * Data = 0.9 * \text{Rate}_{\text{peak}} * T_{\text{position}}$$



$$Data = \frac{0.9 * \text{Rate}_{\text{peak}} * T_{\text{position}}}{1 - 0.9}$$

Maths

We want:

$$\text{Rate}_{\text{effective}} = F \times \text{Rate}_{\text{peak}}$$

$$Data - 0.9 * Data = 0.9 * \text{Rate}_{\text{peak}} * T_{\text{position}}$$

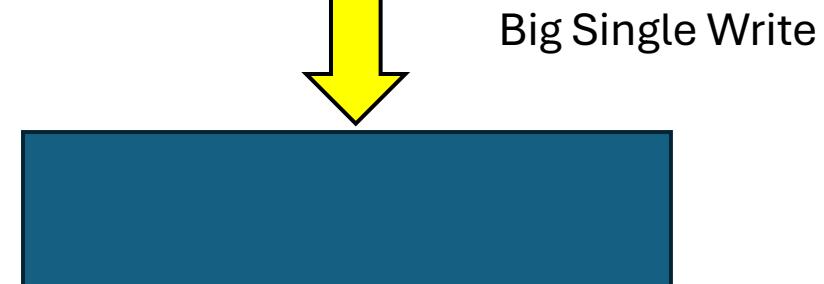
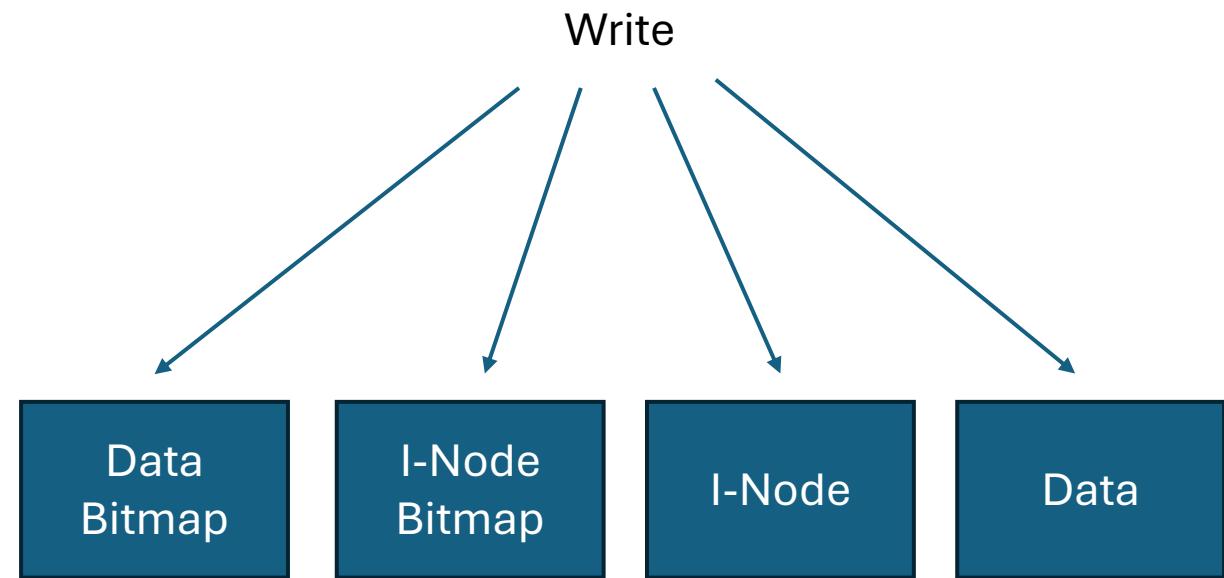
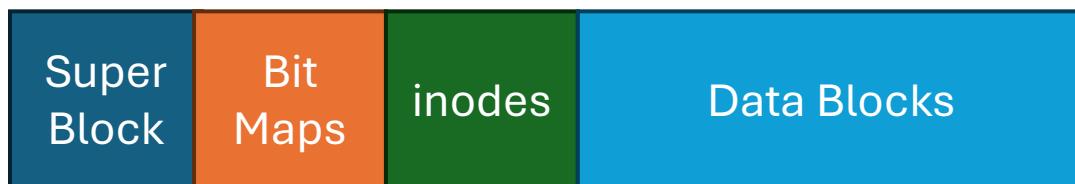

$$Data = \frac{F}{1 - F} * \text{Rate}_{\text{peak}} * T_{\text{position}}$$

We now know how much data we would need to write to get the required level of throughput

Problems

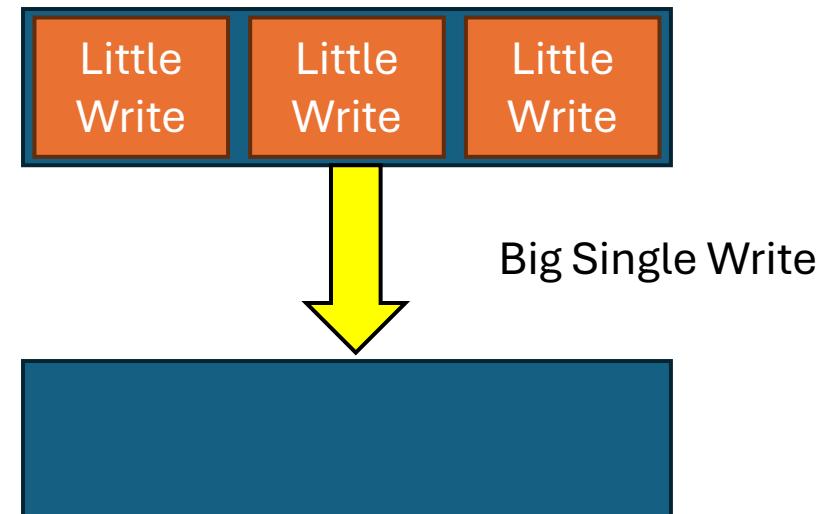
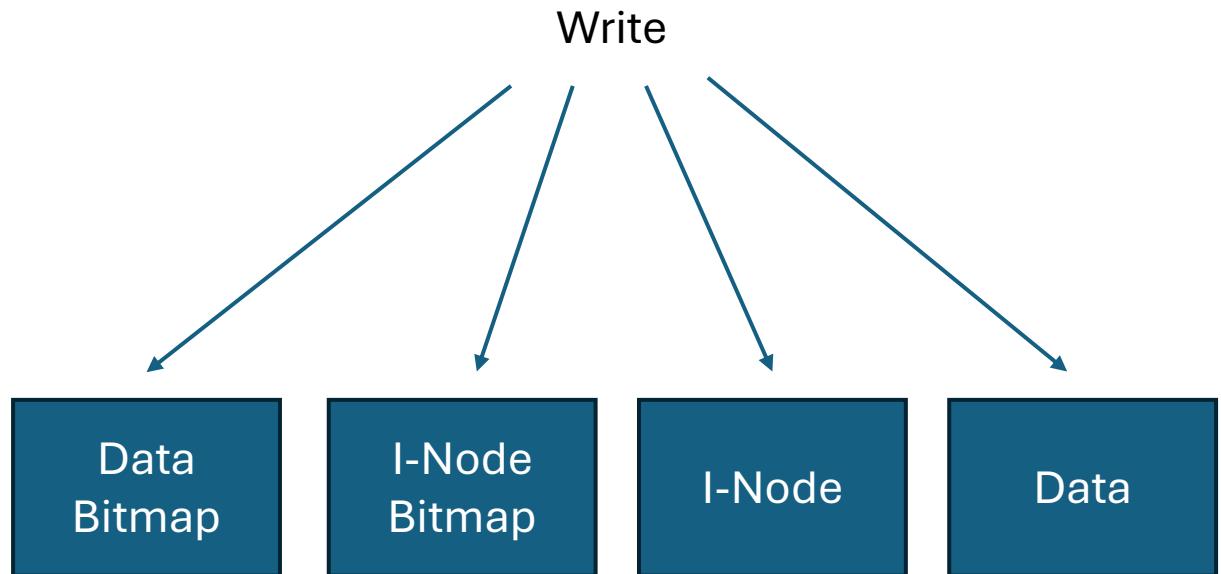
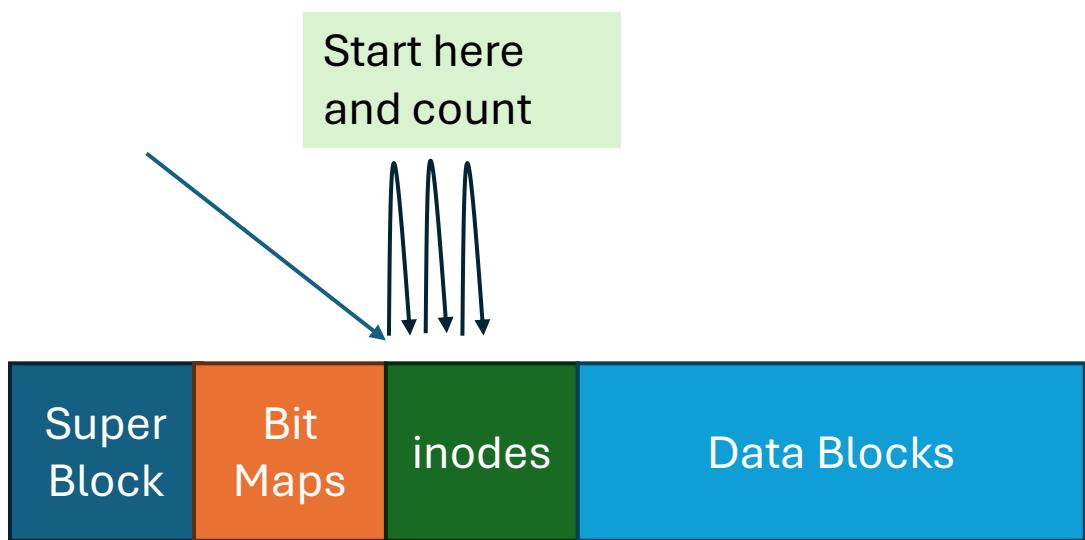
#3: I-Nodes

How we used to do it...



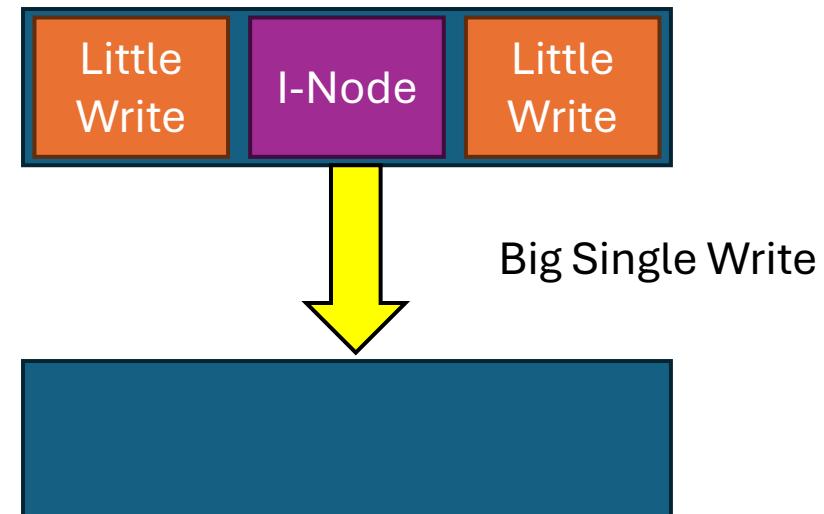
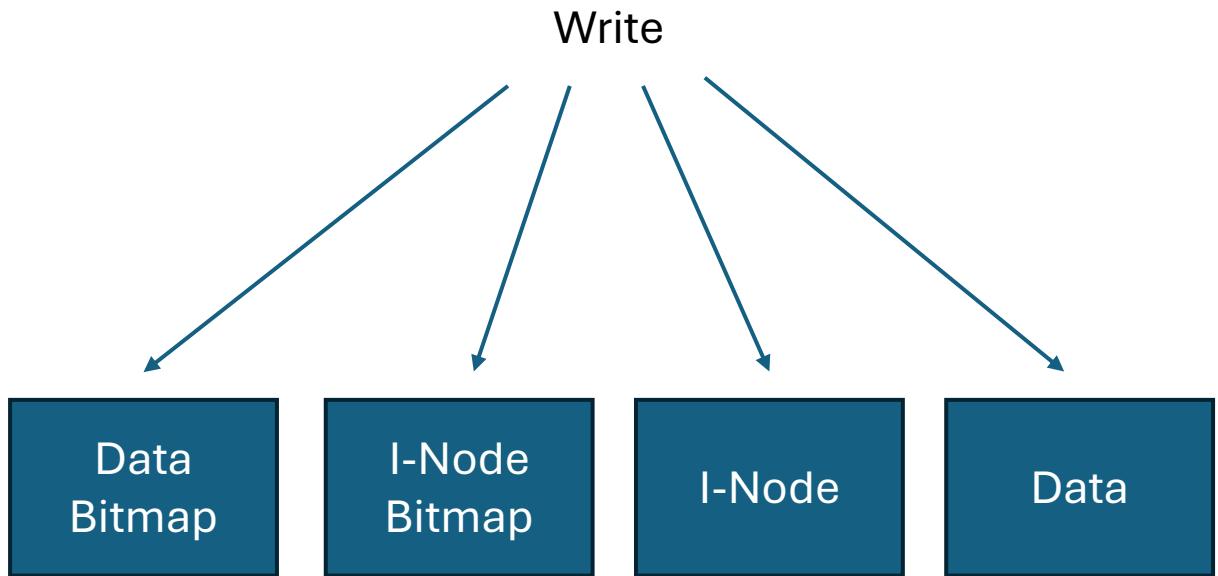
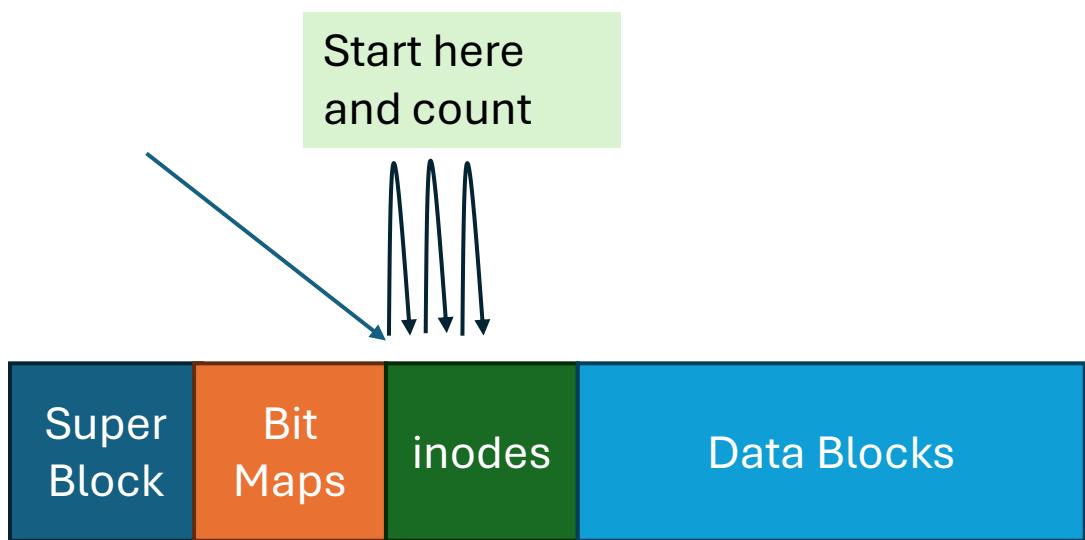
Problems

#3: I-Nodes



Problems

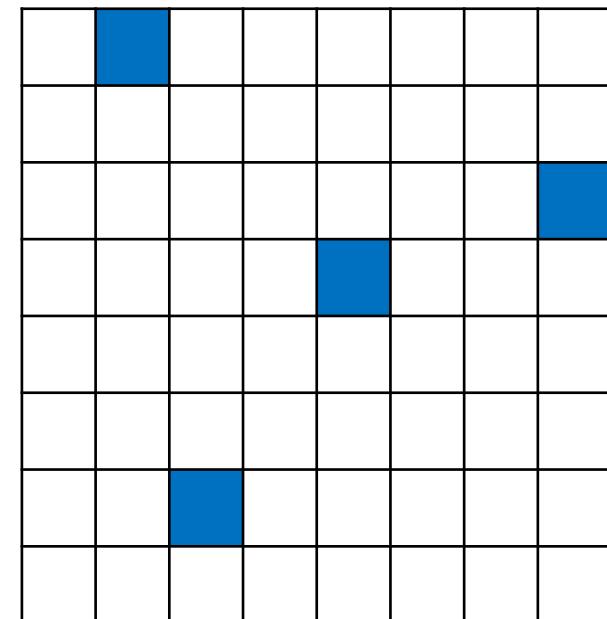
#3: I-Nodes



Problems

#3: Where are the I-Nodes?

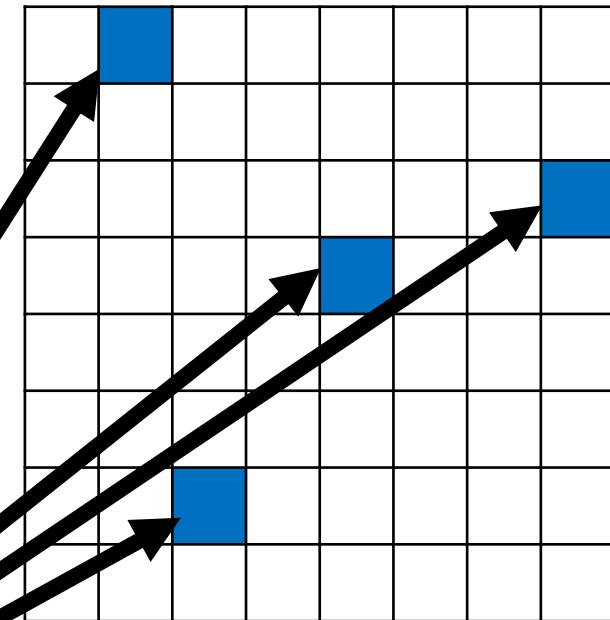
- They are scattered
- The ‘latest version’ keeps moving



Problems

#3: Where are the I-Nodes?

- They are scattered
- The ‘latest version’ keeps moving



We need
a map!

imap

Inode Map (imap)

- Maps inode # to physical address

#	
1	x0000ffff
2	x012e23b2
3	x351f2a21
4	x9eeebac4
...	



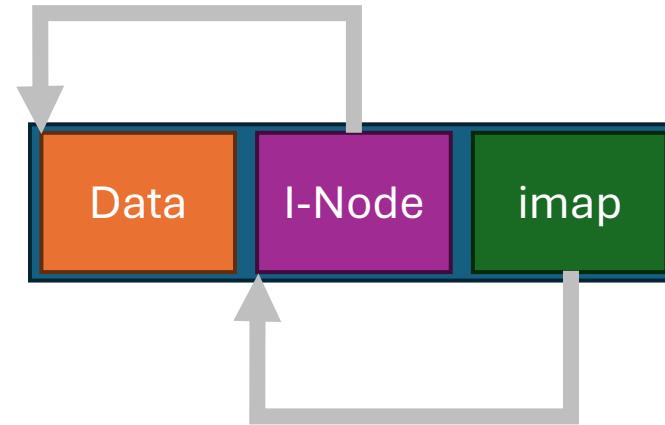
But... where must we
store said map?

In Memory

Writing a Segment

The Cunning Plan

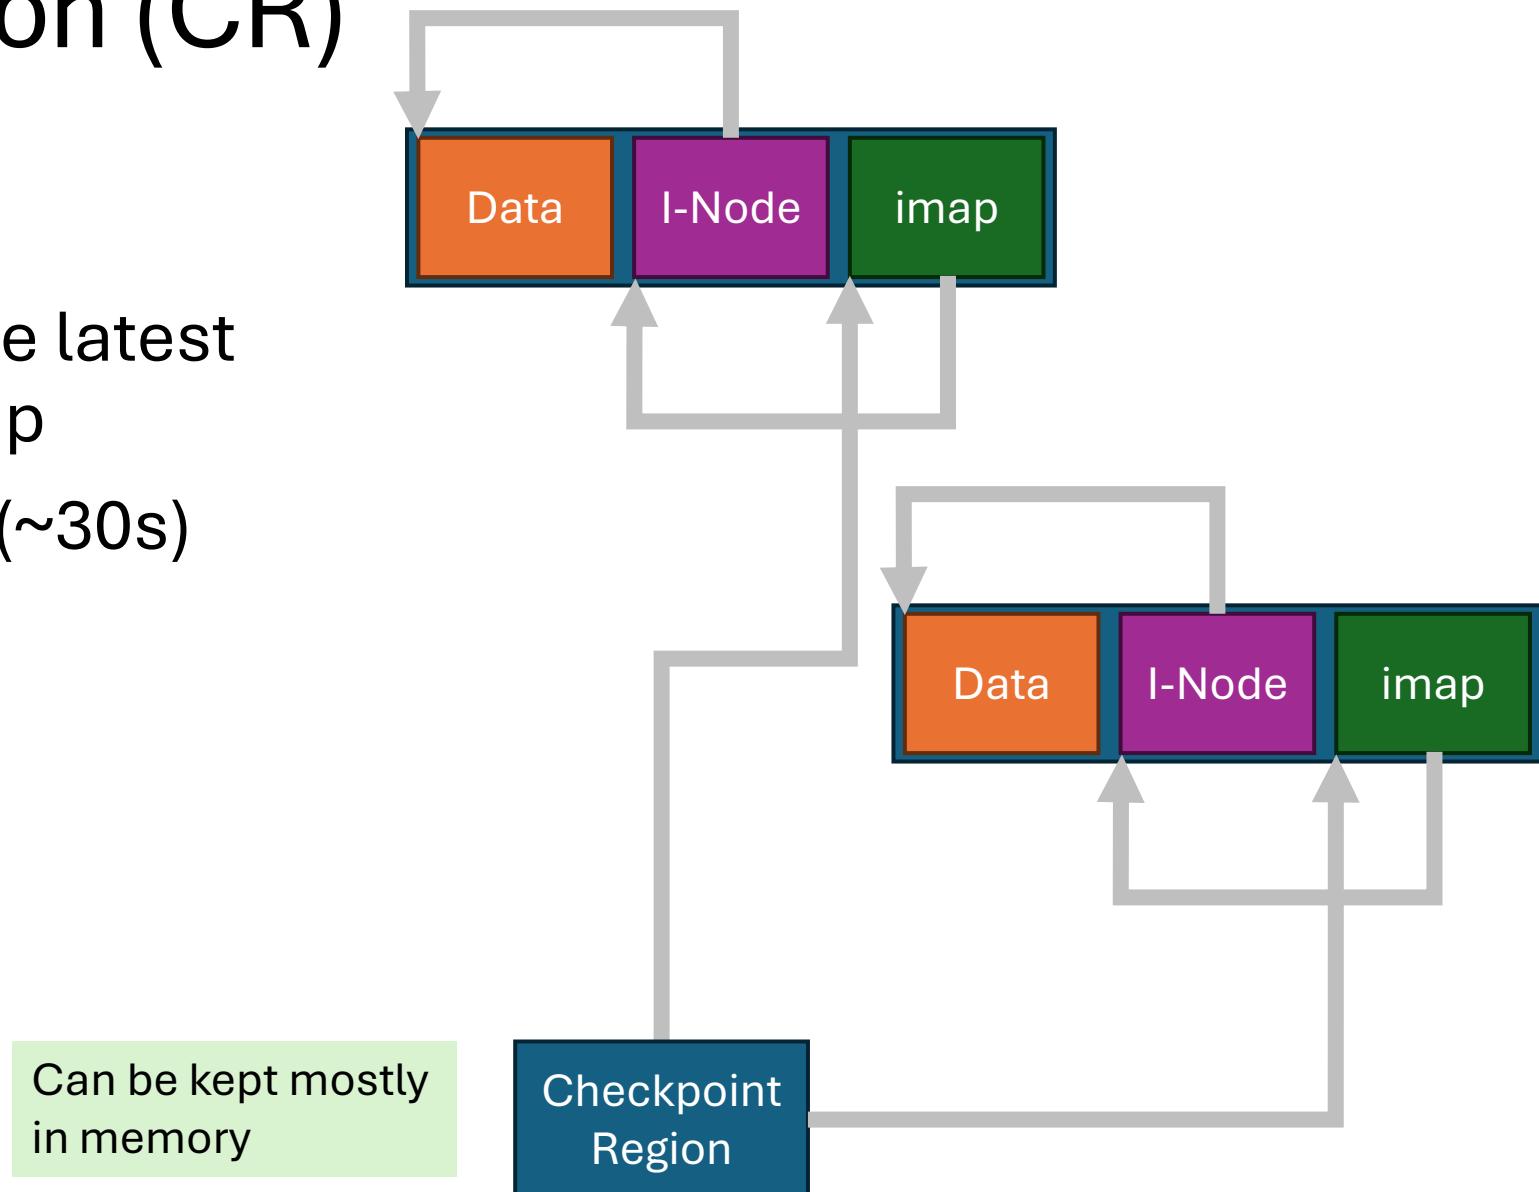
- We write the data
 - But where is the data?
- We write the inode next to the data!
 - But where is the inode?
- We write the imap (or part of it) next to the inode



Checkpoint Region (CR)

Checkpoint Region

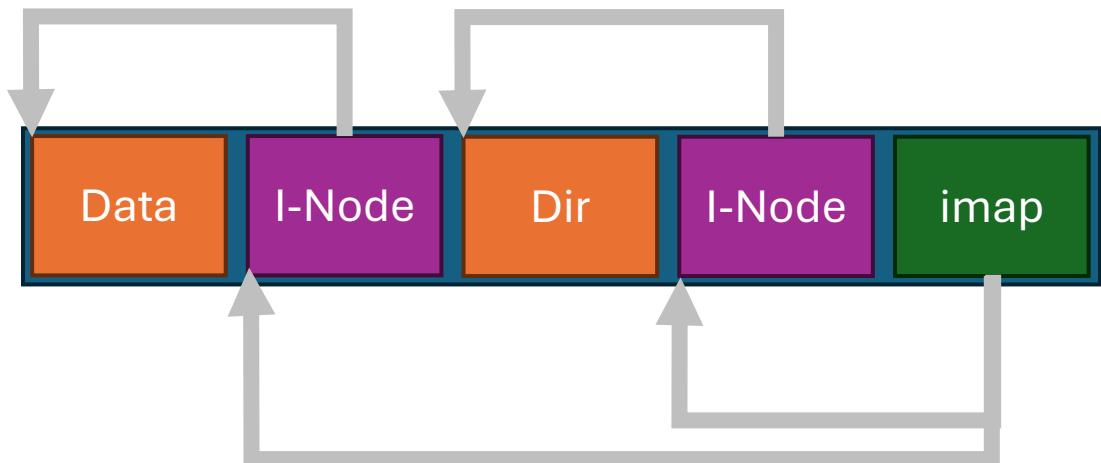
- Contains pointers to the latest pieces of the **inode** map
- Updated in ‘slow time’ (~30s)



Updating Files

Directory Updates

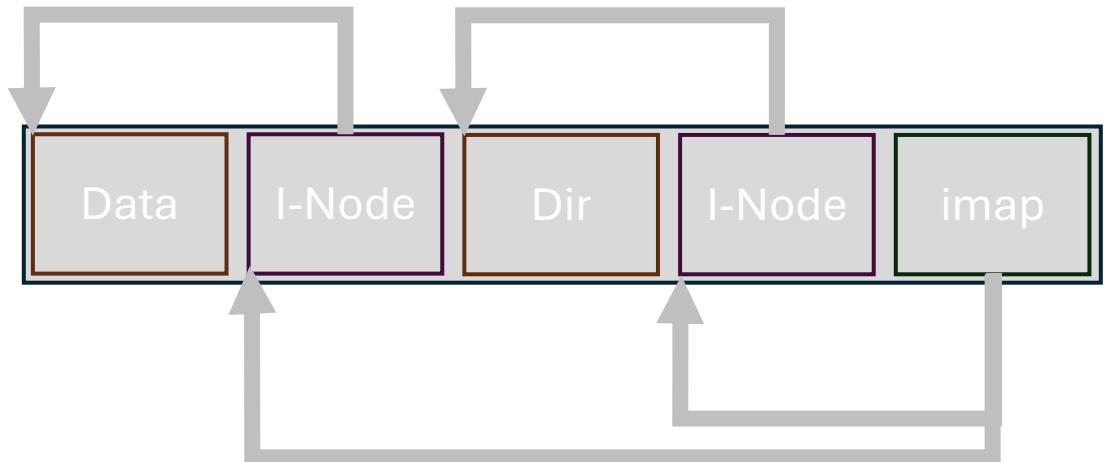
- Creating/modifying a file updates the directory, so...
- This also prevents the **recursive update problem**
(only the inode map is updated with inode locations)



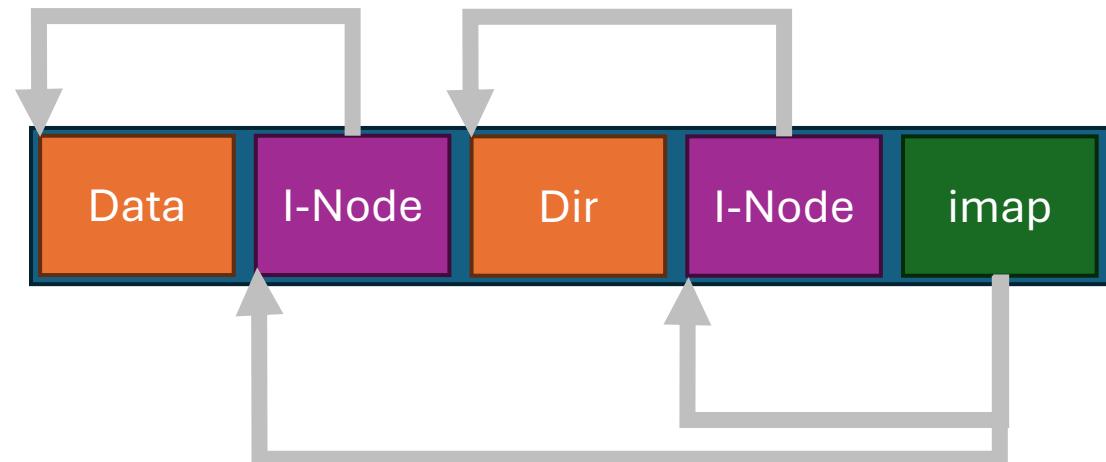
Problems

#4: New ‘writes’ leave old writes...

Old version?



New version?



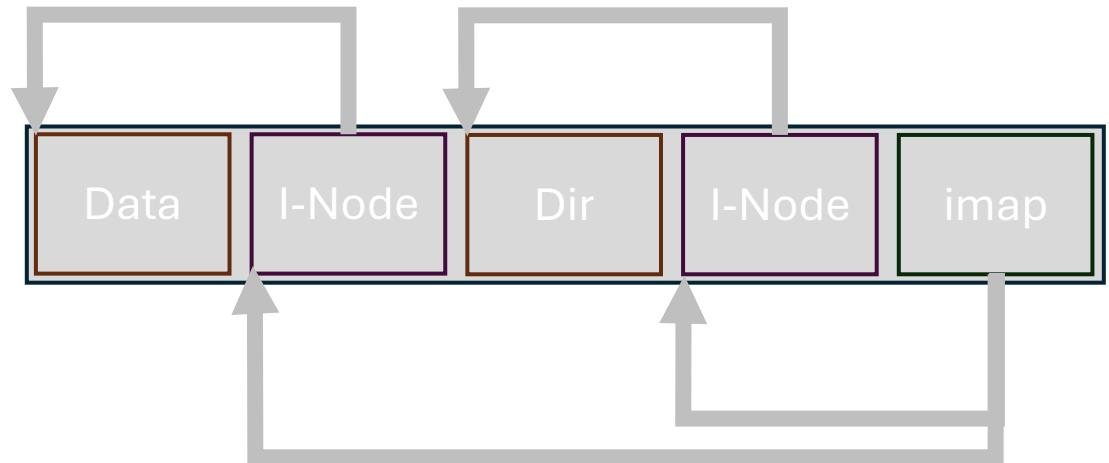
Problems

#4: New ‘writes’ leave old writes...

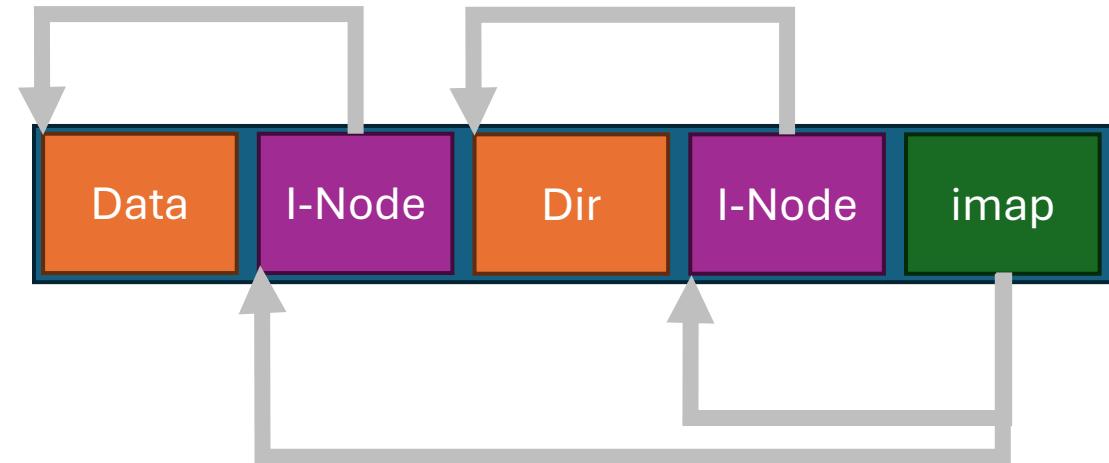
Case 1: Everything new

Case 2: Inode now points to a new block

Old version?



New version?



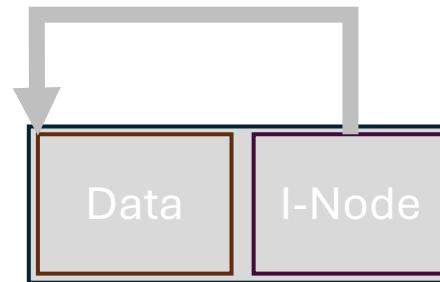
Problems

#4: New ‘writes’ leave old writes...

Case 1: Everything new

Case 2: Inode now points to a new block

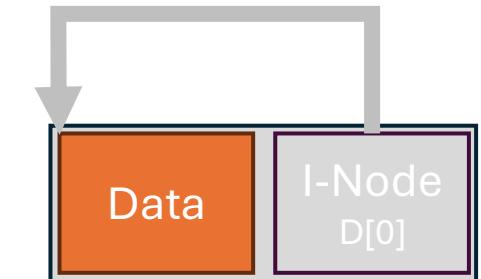
Old version?



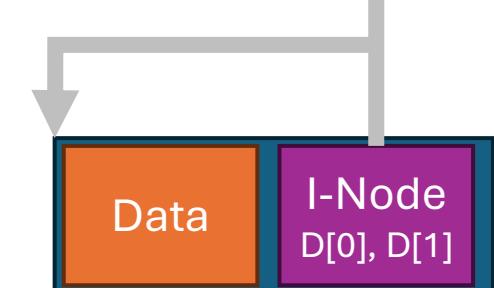
New version?



Old version?



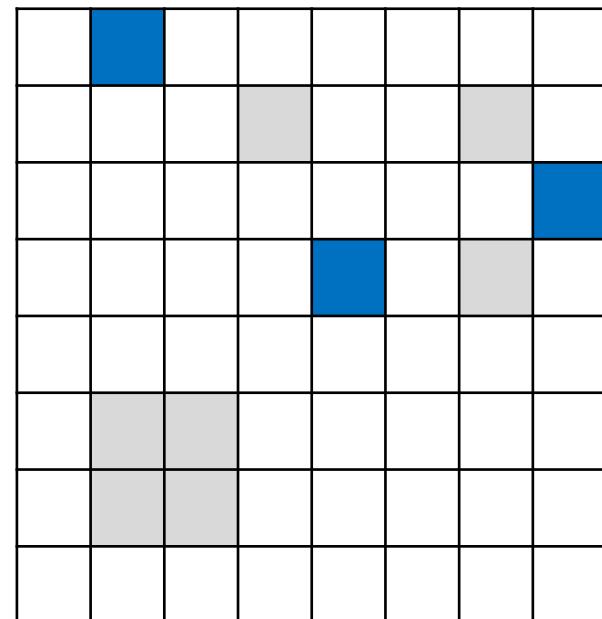
New version?



Problems

#4: New ‘writes’ leave old writes...

Garbage Collection!



Data



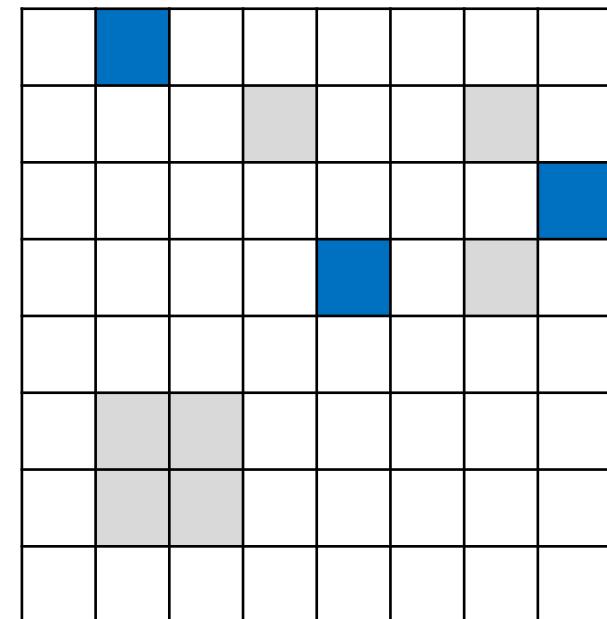
Garbage



Garbage Collection

Problems:

- Need to be careful for holes



Data



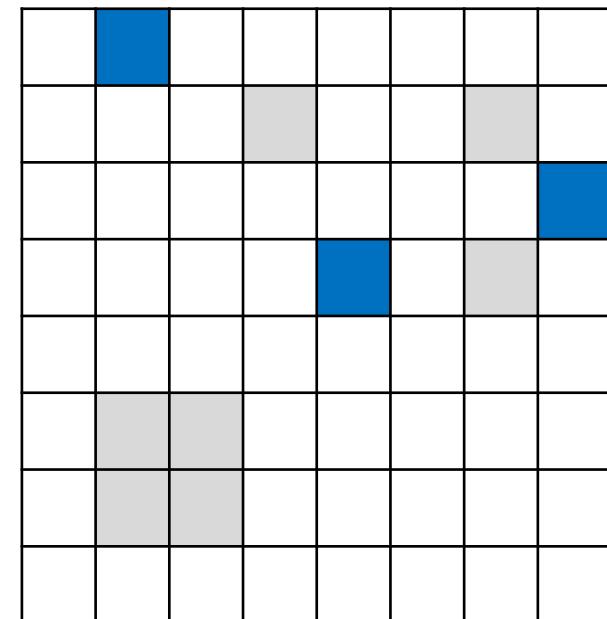
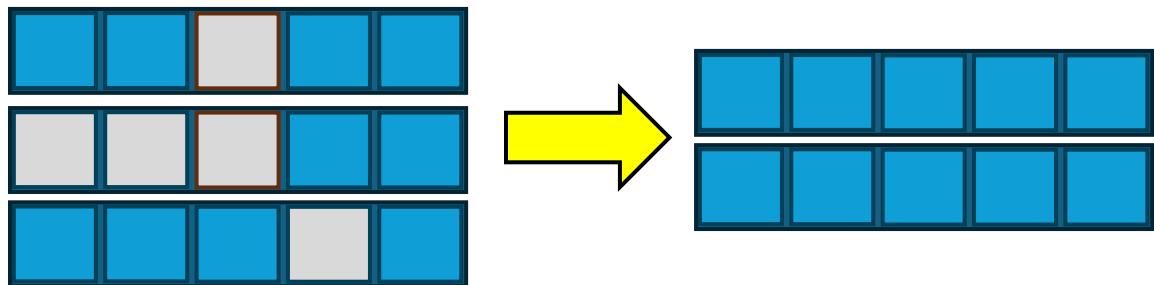
Garbage



Garbage Collection

Problems:

- Need to be careful for holes



Data



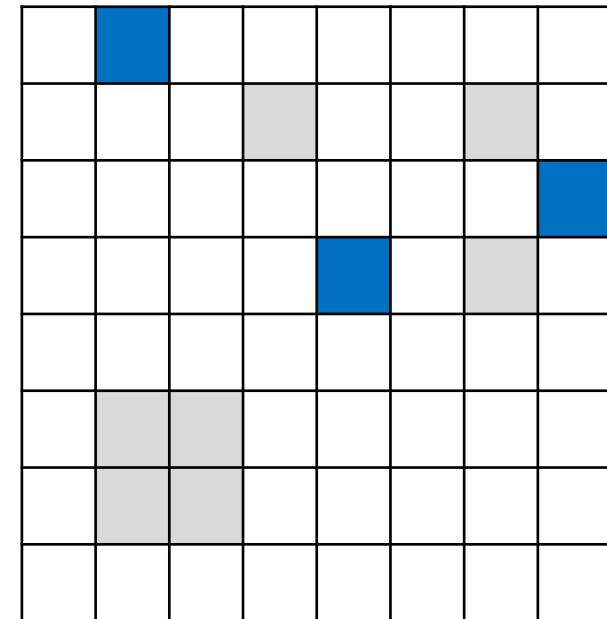
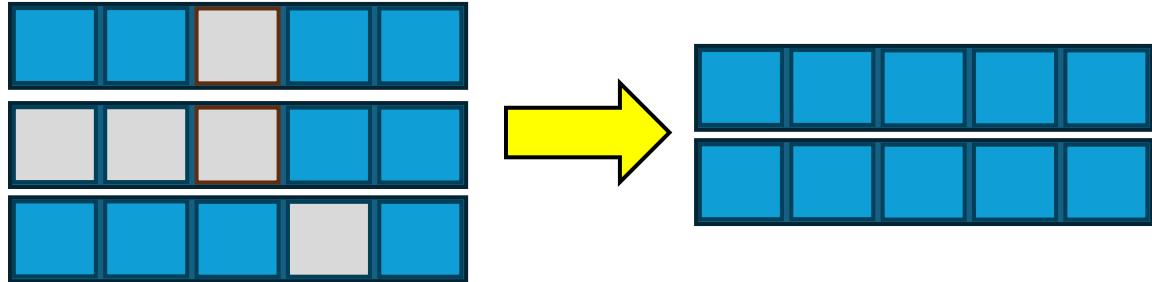
Garbage



Garbage Collection

Questions:

- How to track ‘dead blocks’?
- How often to collect garbage?



Data



Garbage



Block Liveness

- Each segment has a **segment summary block**

Process

- Read **segment summary block**
- Find **inode number**
- Check **imap** (follow to the data)
- Is it where it says it should be?

Segment Summary Block

- Data Block:
 - inode number
 - offset (which block am I?)



What and When?

When?

- Periodically
- When you're bored (idle time)

What?

- Subject to research...

Thoughts:

- Frequently changing files are collected less frequently?
- Infrequently changing files are collected more frequently?

Summary

- LFS: Log-Structured File System

Problems

#5: Crash...

- **While writing a segment**
- While writing the checkpoint region

We also do some **logging...**

Problems

#5: Crash...

- While writing a segment
- **While writing the checkpoint region**

We have two... one at each end with alternate writing...

We can use timestamps and consistency to work out which one is correct (basically versioning)

Reminder, the **checkpoint region** is where we store the references to the imap chunks

Problems

#5: Crash...

- While writing a segment
- **While writing the checkpoint region**

We have two... one at each end with alternate writing...

We can use timestamps and consistency to work out which one is correct (basically versioning)

Reminder, the **checkpoint region** is where we store the references to the imap chunks

Problems

#5: Crash...

- While writing a segment
- **While writing the checkpoint region**

We lose ~30s of requests

- Can fix through **roll forward**:
- Read through the log finding valid updates...

Questions?

